

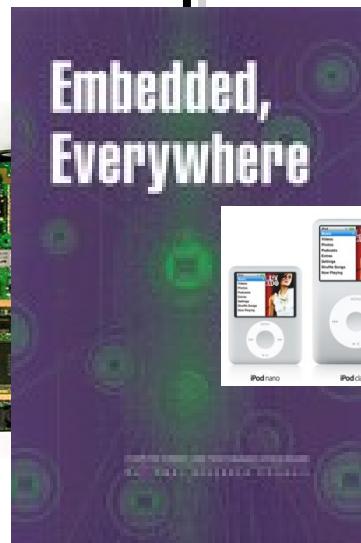
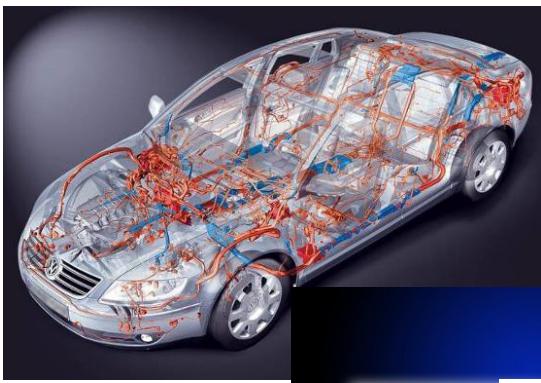
# ARM System Introduction

Jim Huang <jserv@0xlab.org>

September 25, 2012

Modifications from Prabal Dutta, University of Michigan

# Embedded, everywhere



What is driving the  
embedded everywhere explosion?

# Moore's Law: IC transistor count doubles every two years

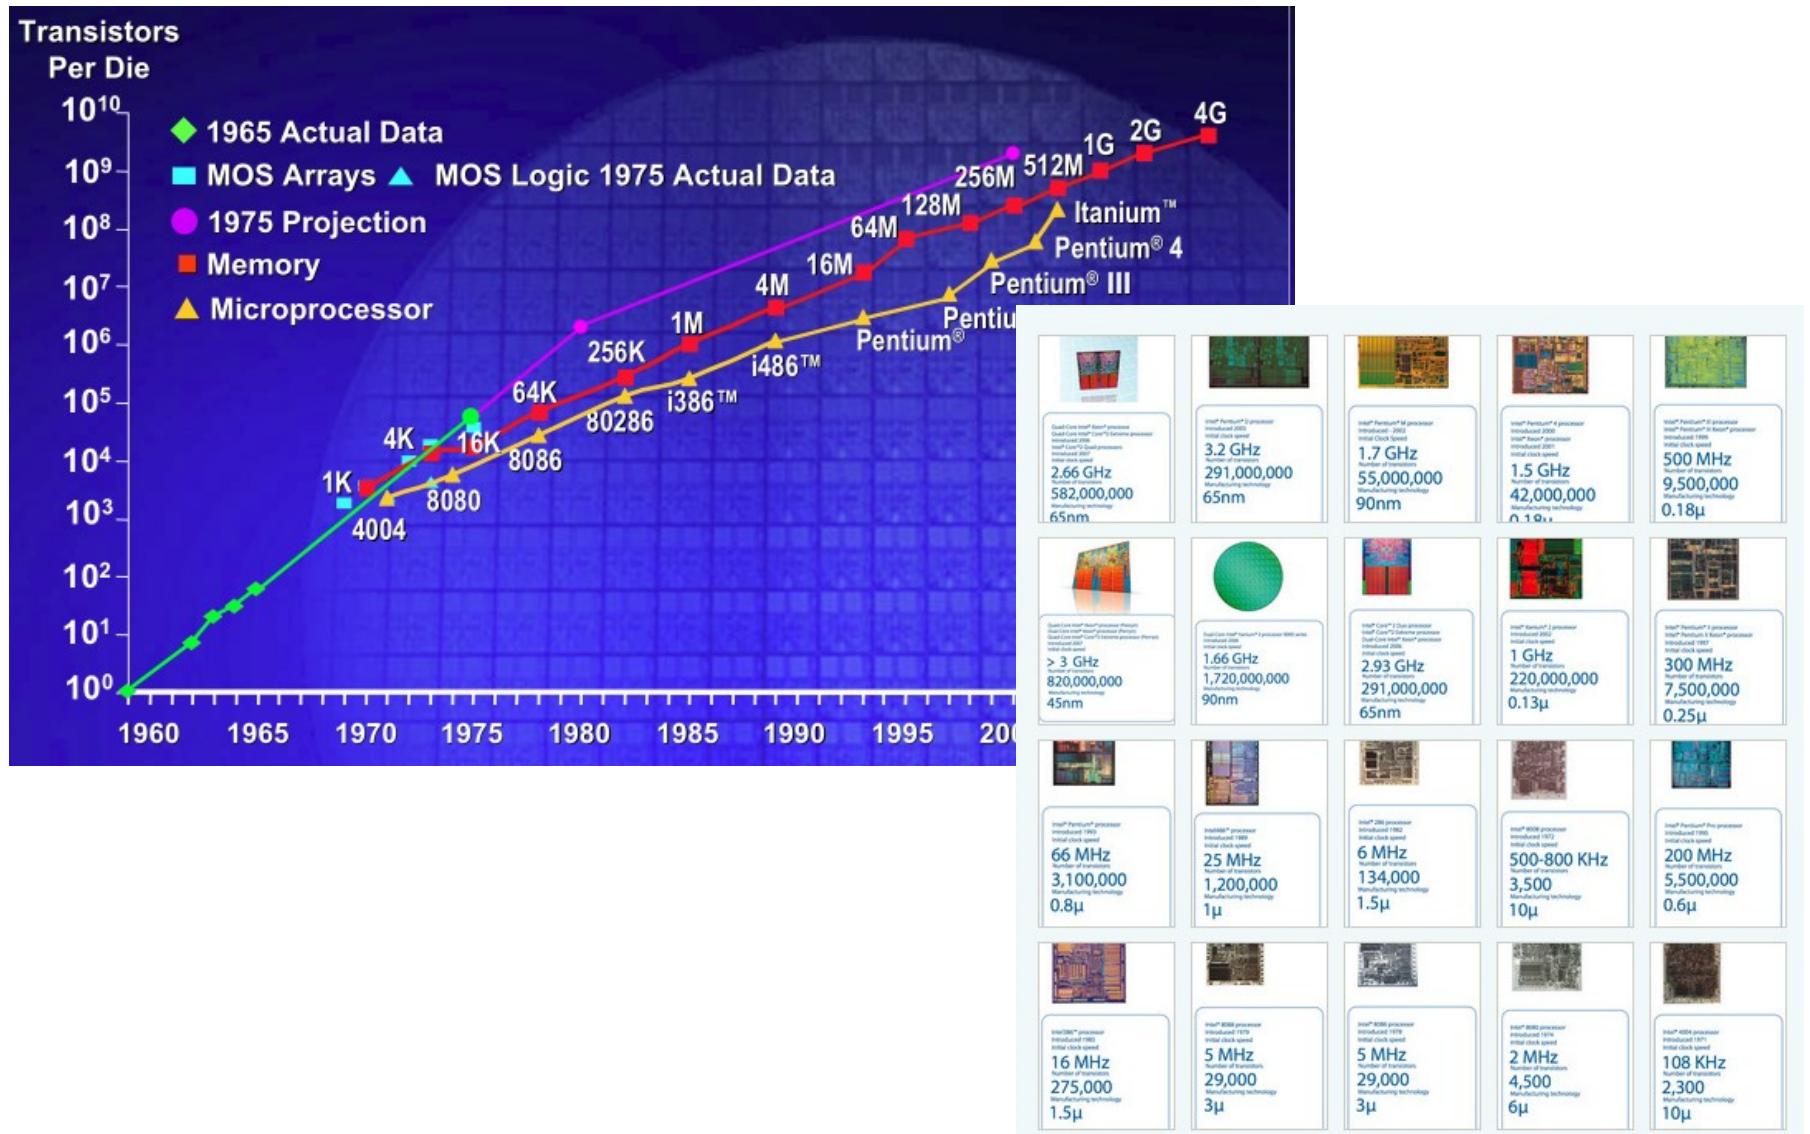


Photo Credit: Intel

# Flash memory scaling: Rise of density & volumes; Fall (and rise) of prices

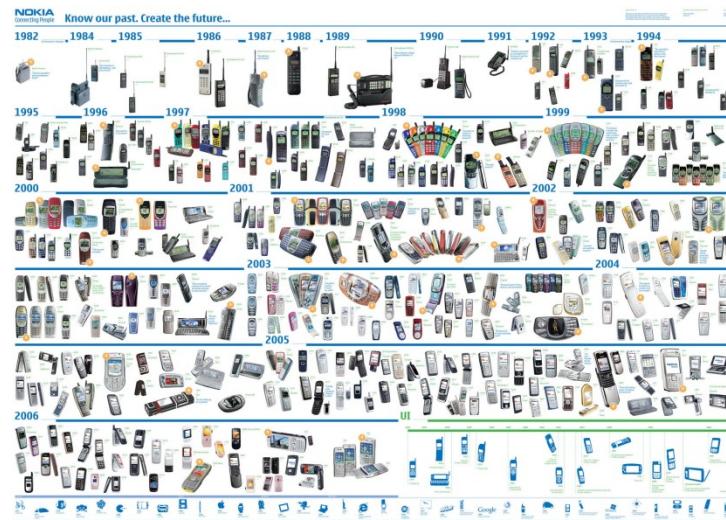
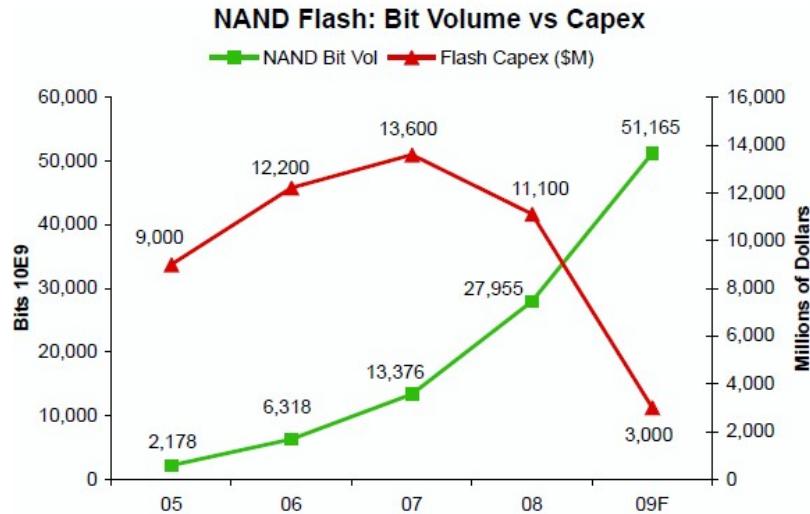
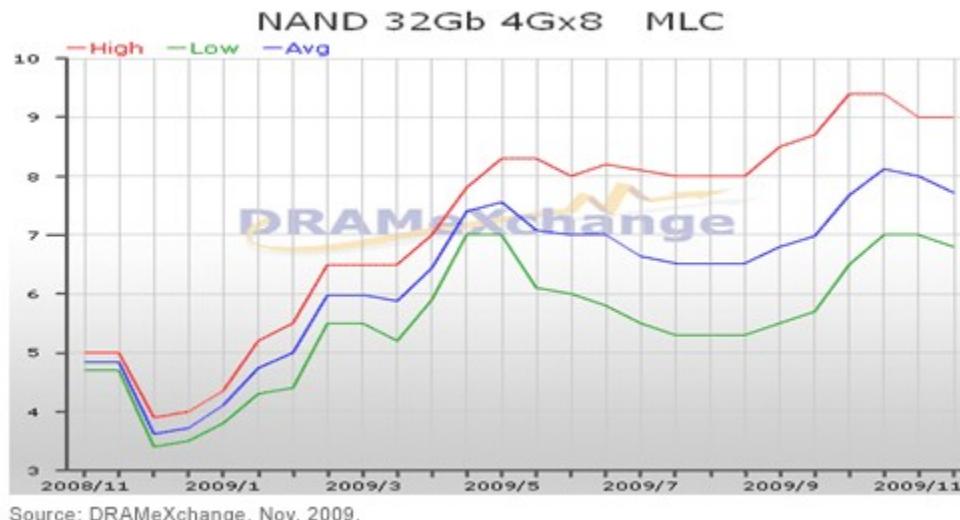
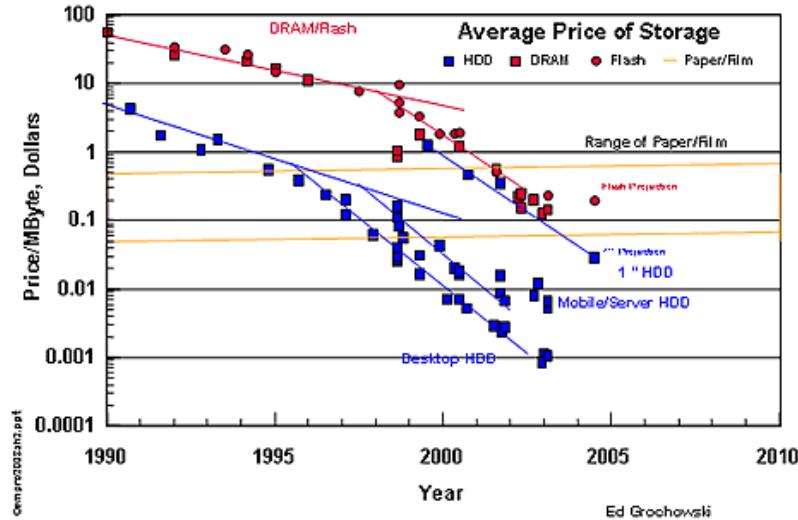
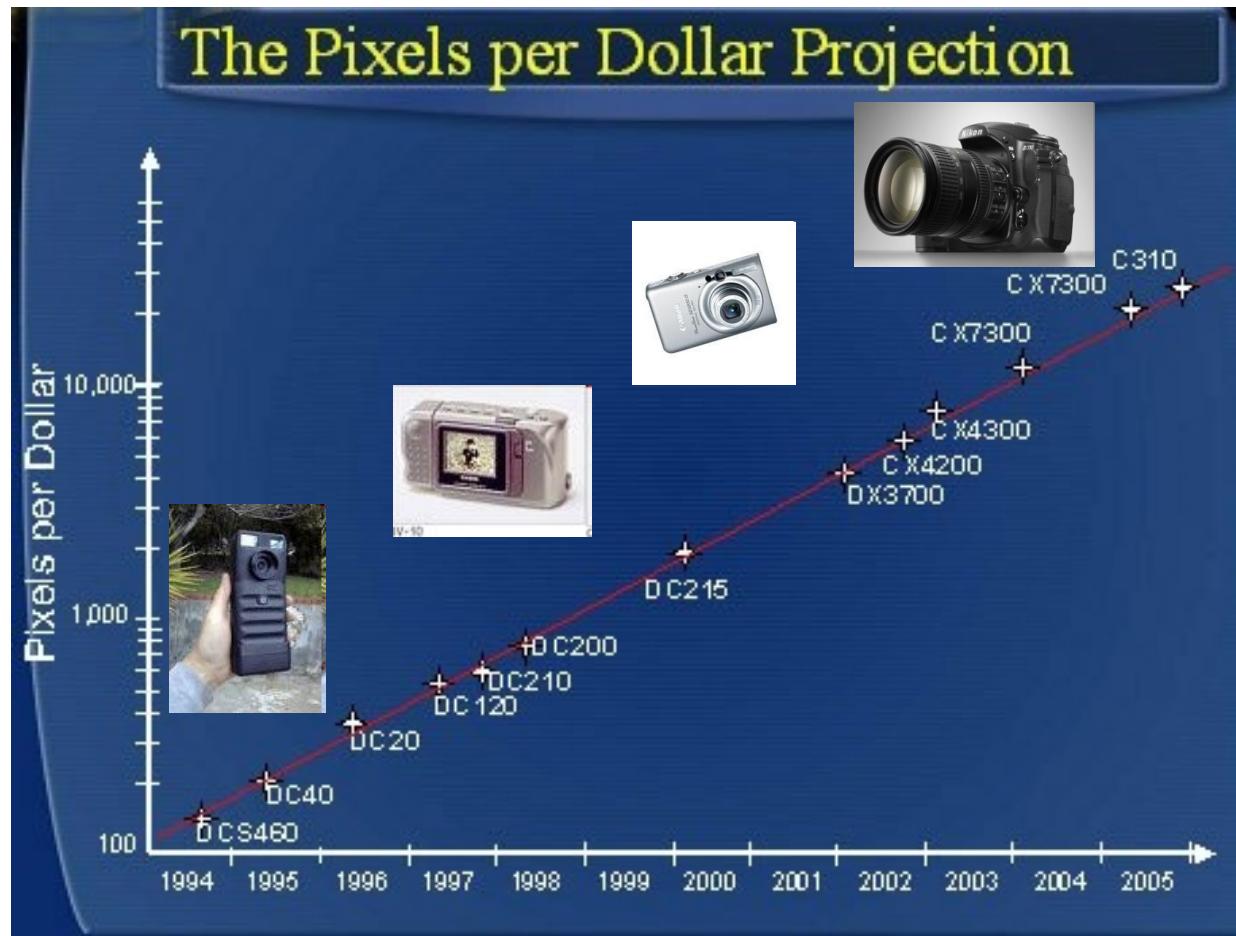


Figure-1 32Gb MLC NAND Flash contract price trend

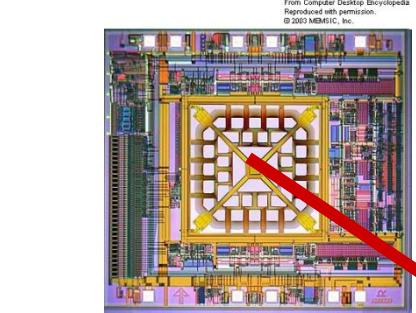


# Hendy's “Law”: Pixels per dollar doubles annually

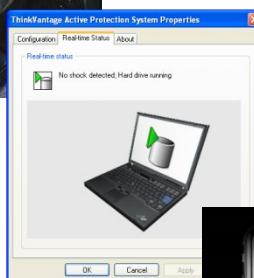


Credit: Barry Hendy/Wikipedia

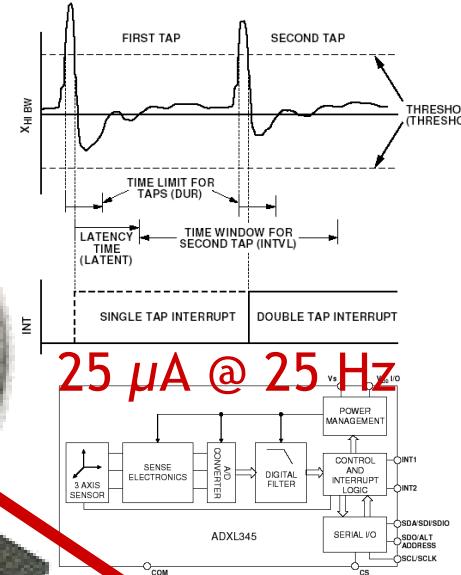
# MEMS Accelerometers: Rapidly falling price and power



O(mA)



ADXL345  
[Analog Devices, 2009]

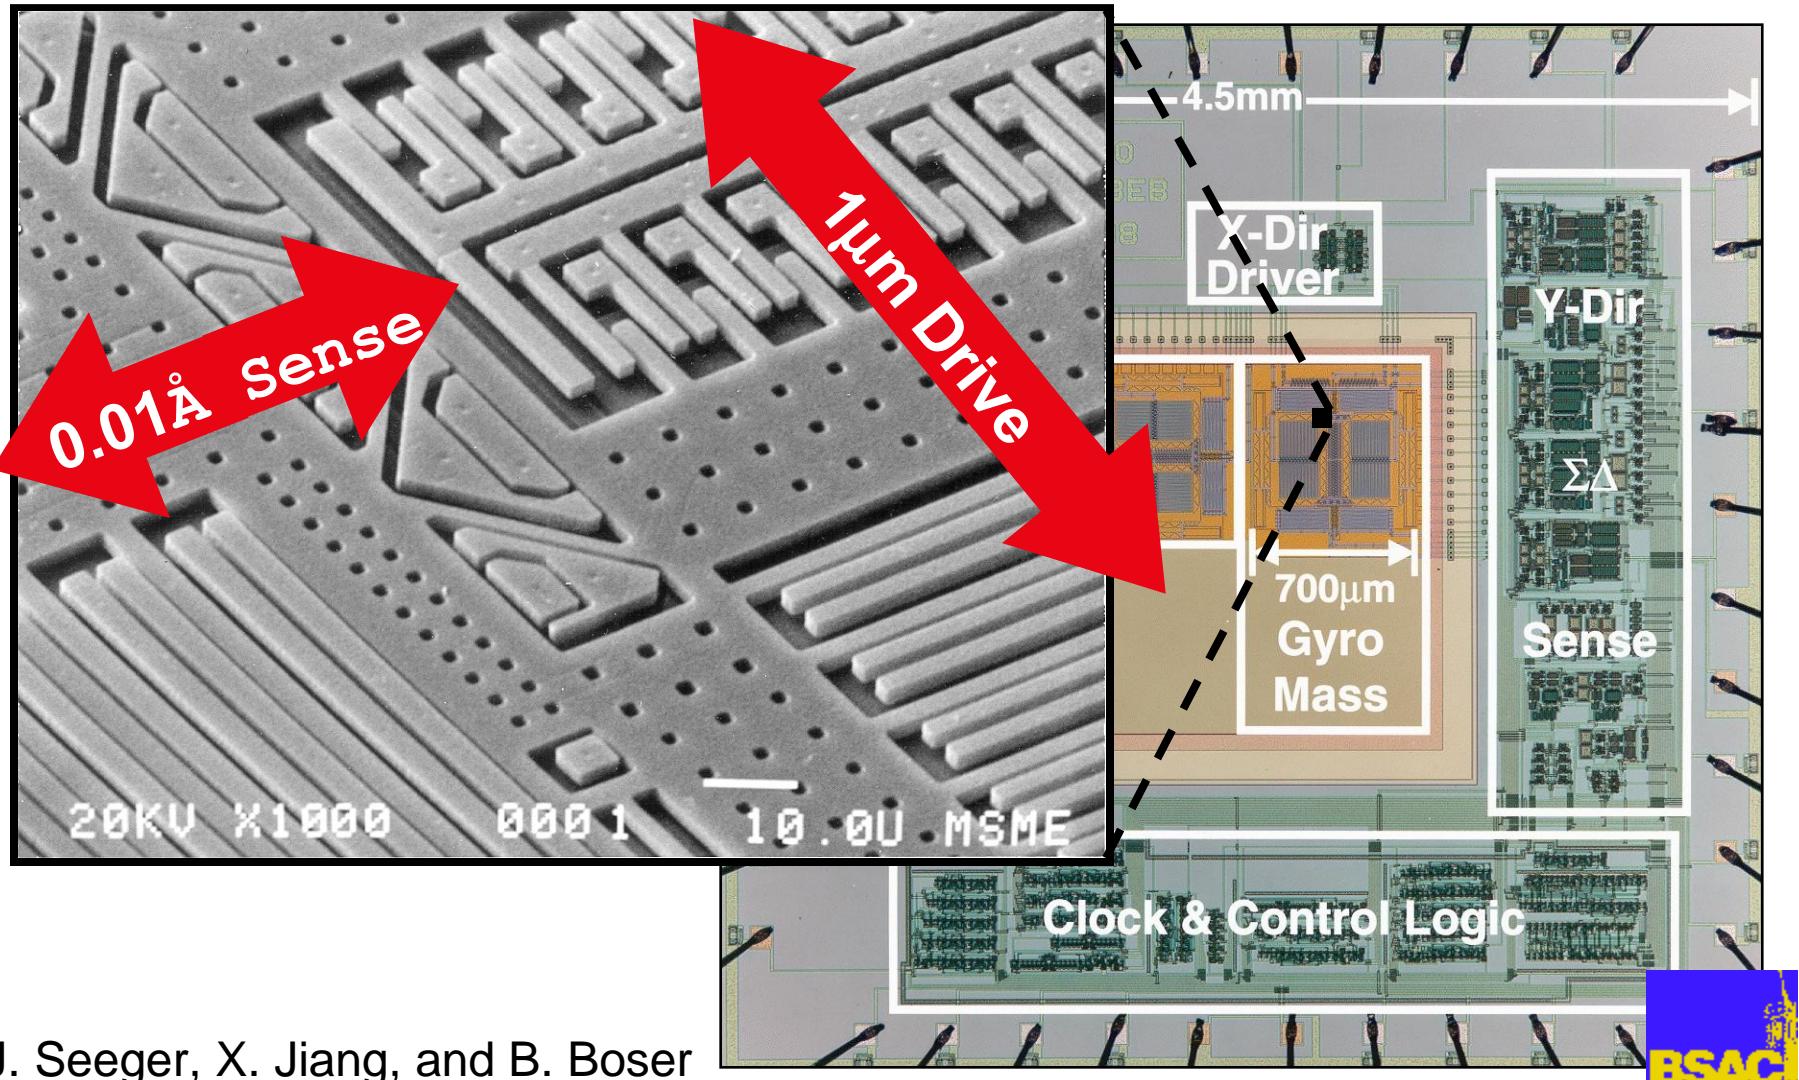


25  $\mu$ A @ 25 Hz



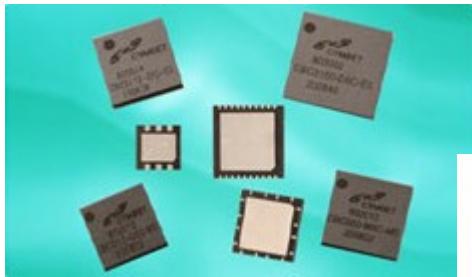
10  $\mu$ A @ 10 Hz @ 6 bits  
[ST Microelectronics, ann. 2009]

# MEMS Gyroscope Chip

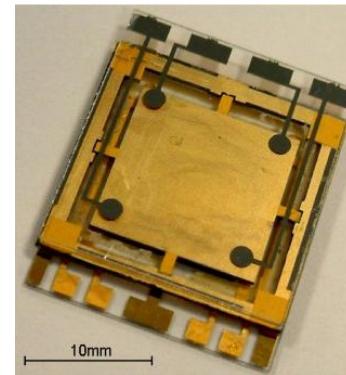


J. Seeger, X. Jiang, and B. Boser

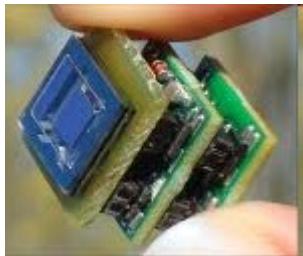
# Energy harvesting and storage: Small doesn't mean powerless...



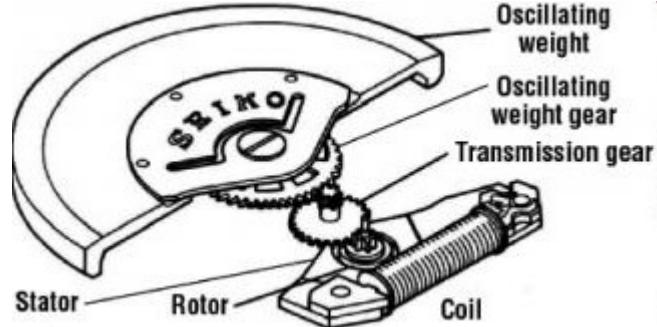
Thin-film batteries



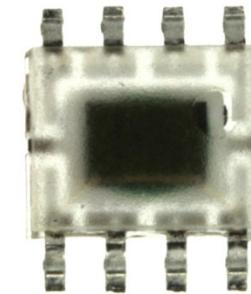
Electrostatic Energy  
Harvester [ICL]



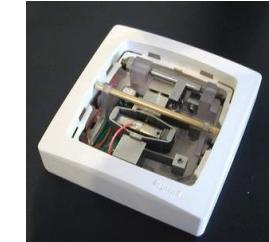
Piezoelectric  
[Holst/IMEC]



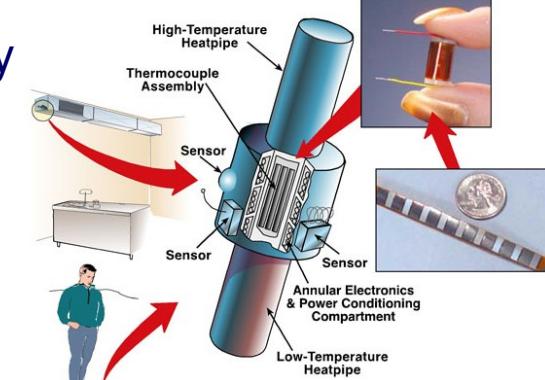
RF [Intel]



Clare Solar Cell

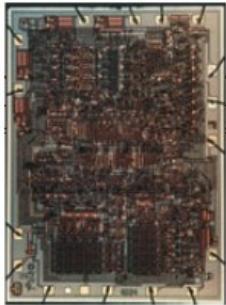


Shock Energy Harvesting  
CEDRAT Technologies



Thermoelectric Ambient  
Energy Harvester [PNNL]

# Bell's Law, Take 2: Corollary to the Laws of Scale



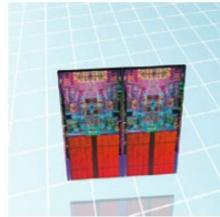
Intel® 4004 processor  
Introduced 1971  
Initial clock speed

108 KHz  
Number of transistors

2,300  
Manufacturing technology

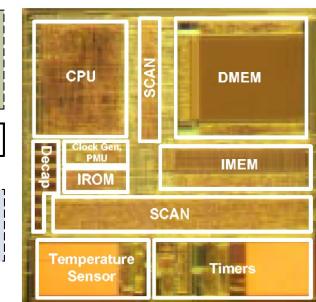
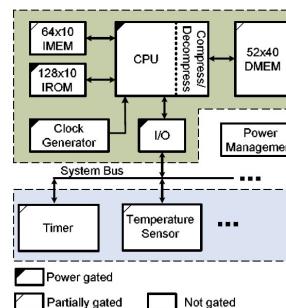
10 $\mu$

15x size decrease  
40x transistors  
55x smaller  $\lambda$



Quad-Core Intel® Xeon® processor  
Quad-Core Intel® Core™2 Extreme processor  
Introduced 2006  
Intel® Core™2 Quad processors  
Introduced 2007  
Initial clock speed

2.66 GHz  
Number of transistors  
582,000,000  
Manufacturing technology  
65nm



UMich Phoenix Processor

Introduced 2008

Initial clock speed

106 kHz @ 0.5V Vdd

Number of transistors

92,499

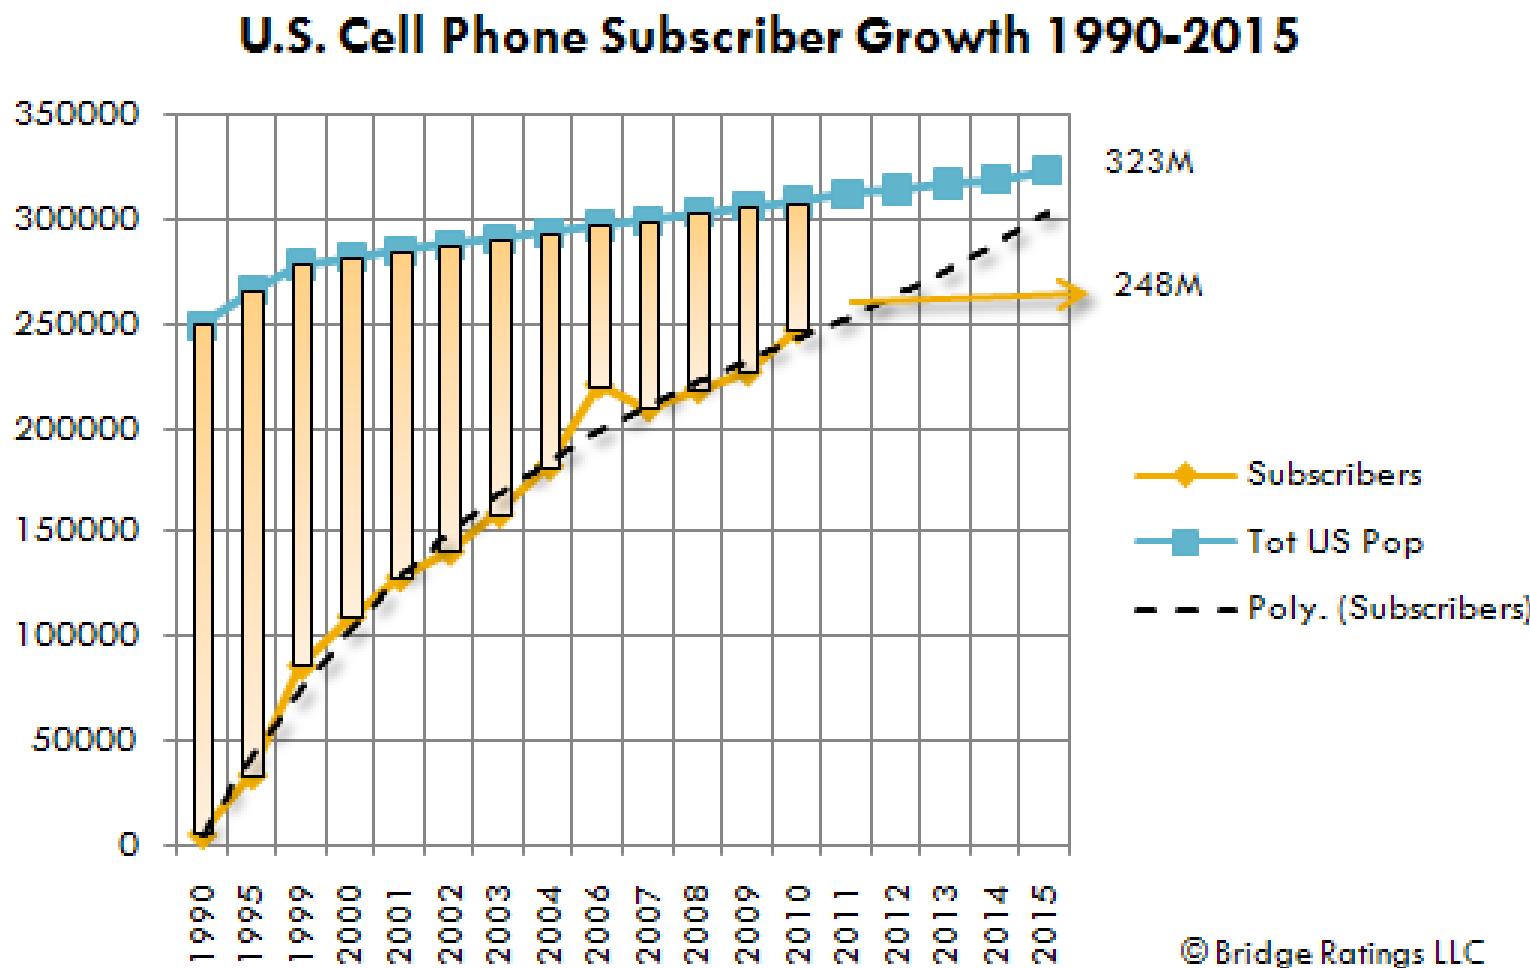
Manufacturing technology

0.18  $\mu$

Photo credits: Intel, U. Michigan

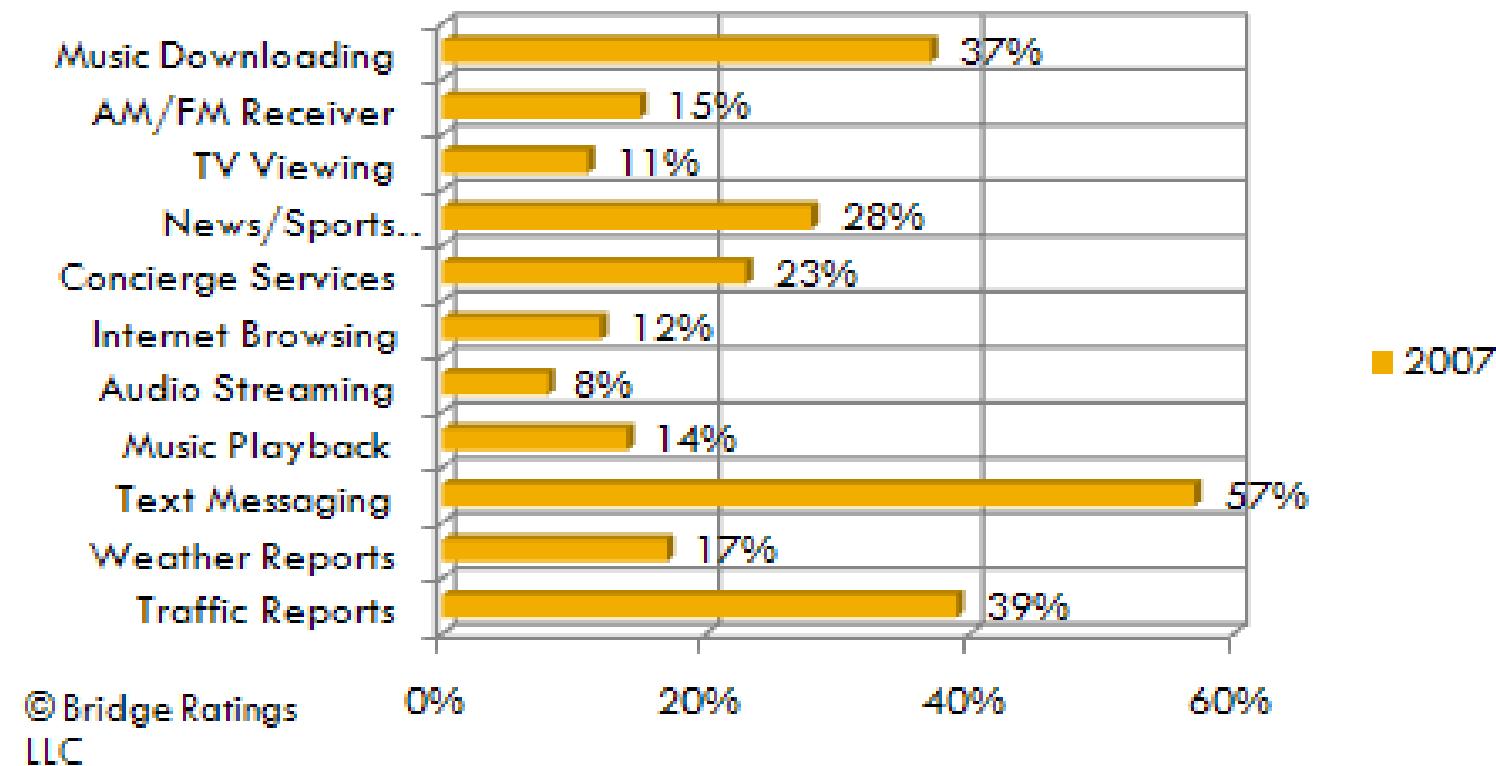
Learning happens when assumptions  
are challenged and invalidated, so...

# Mobile phones: the most successful technology ever



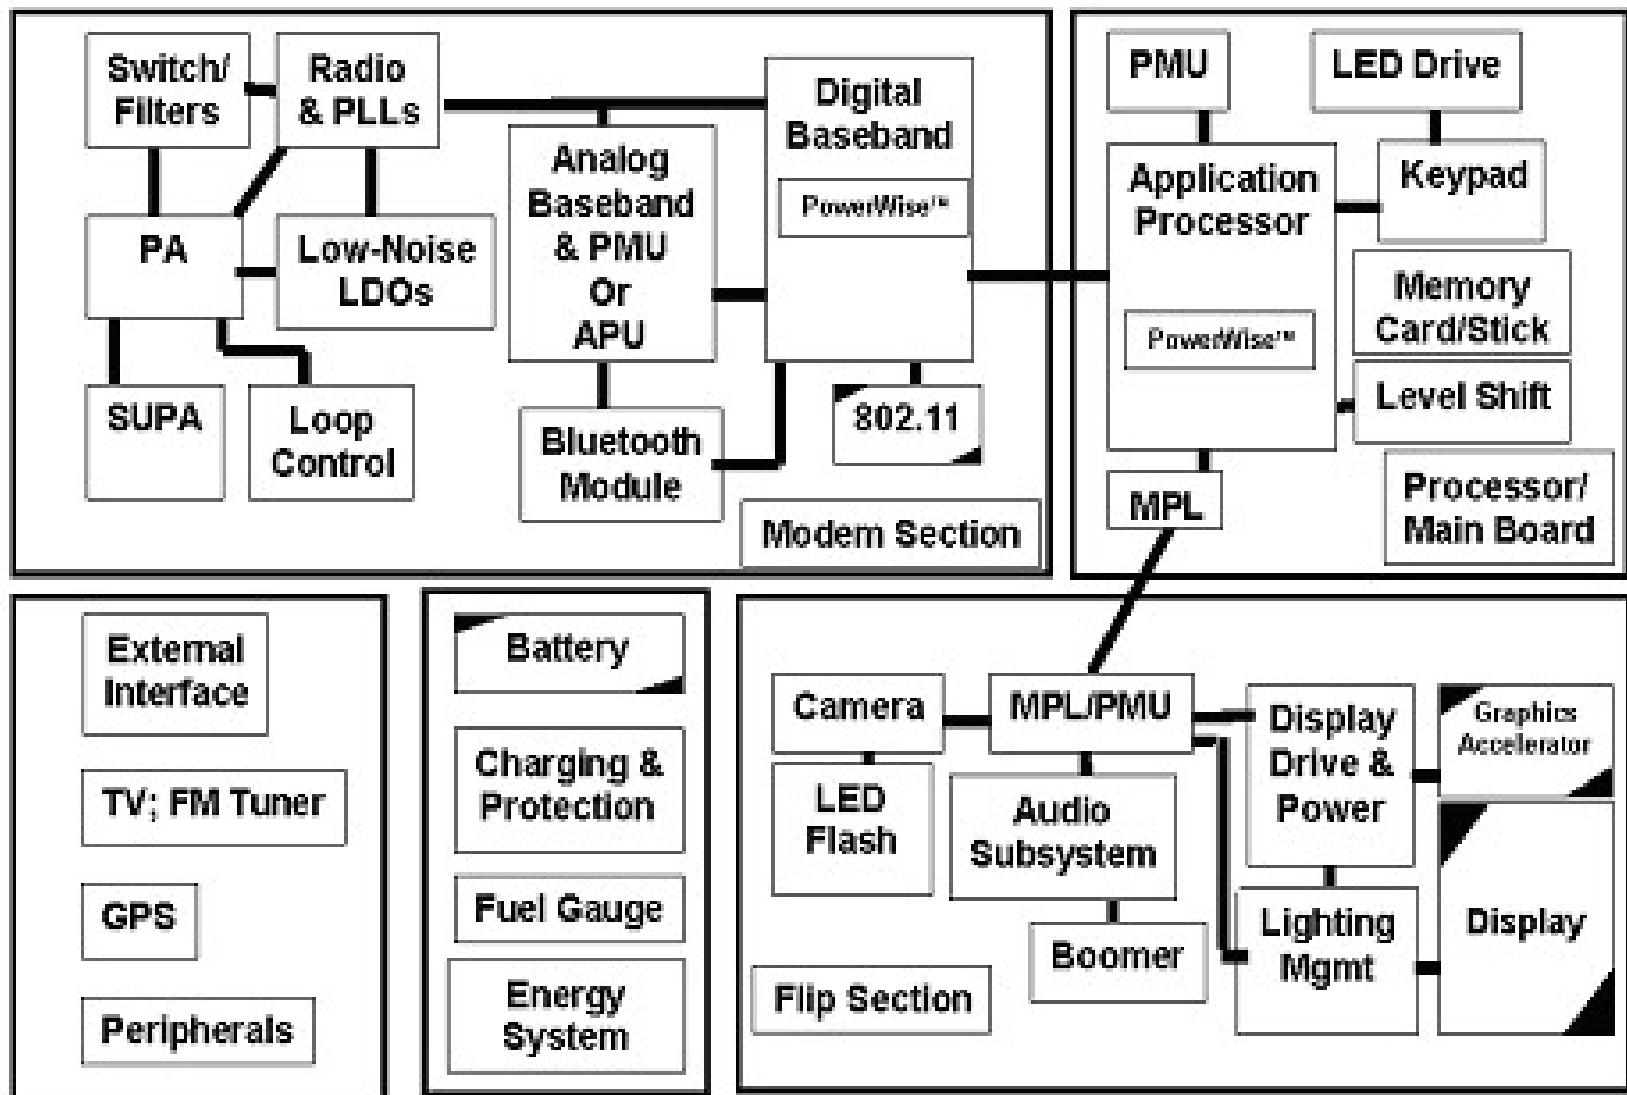
# What happened elsewhere now happens on the phone

## Preferred Cell Phone Services



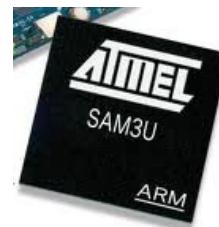
What happens when you press the power switch on your mobile phone?

# Mobile phone system architecture



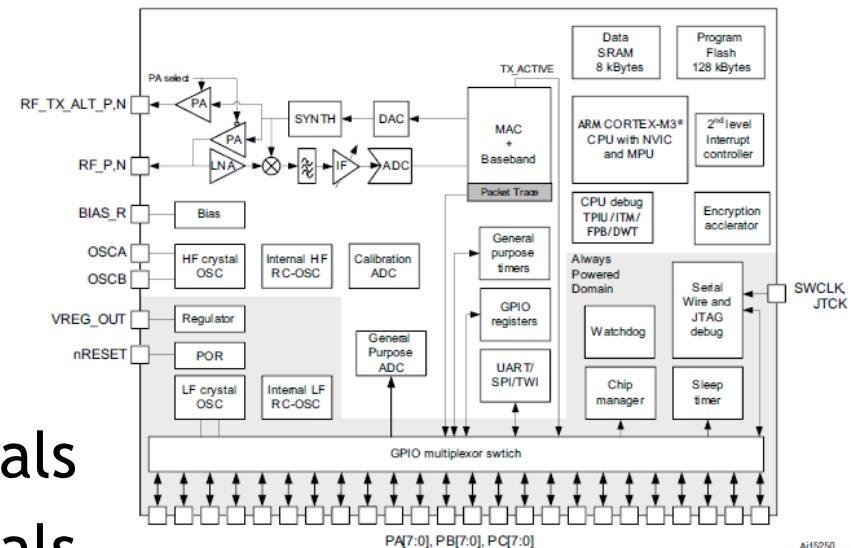
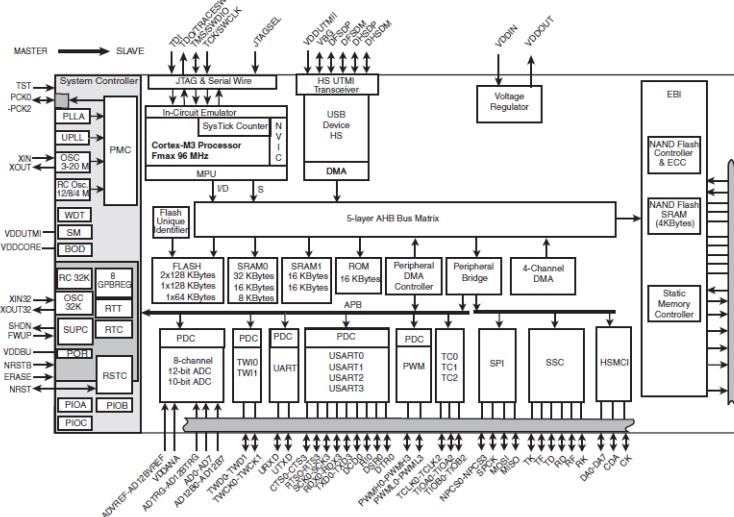
And, most of them are ARM powered

# Lots of manufacturers ship ARM products

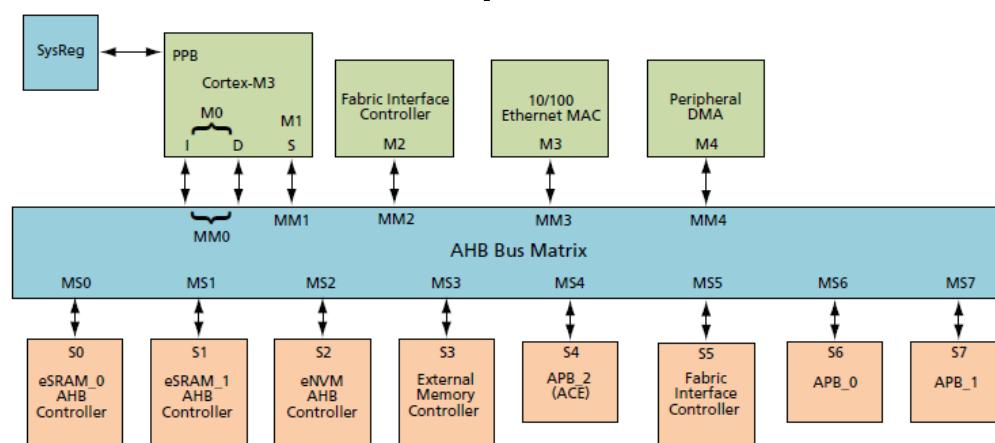


What differentiates these  
products from one another?

# The difference is...

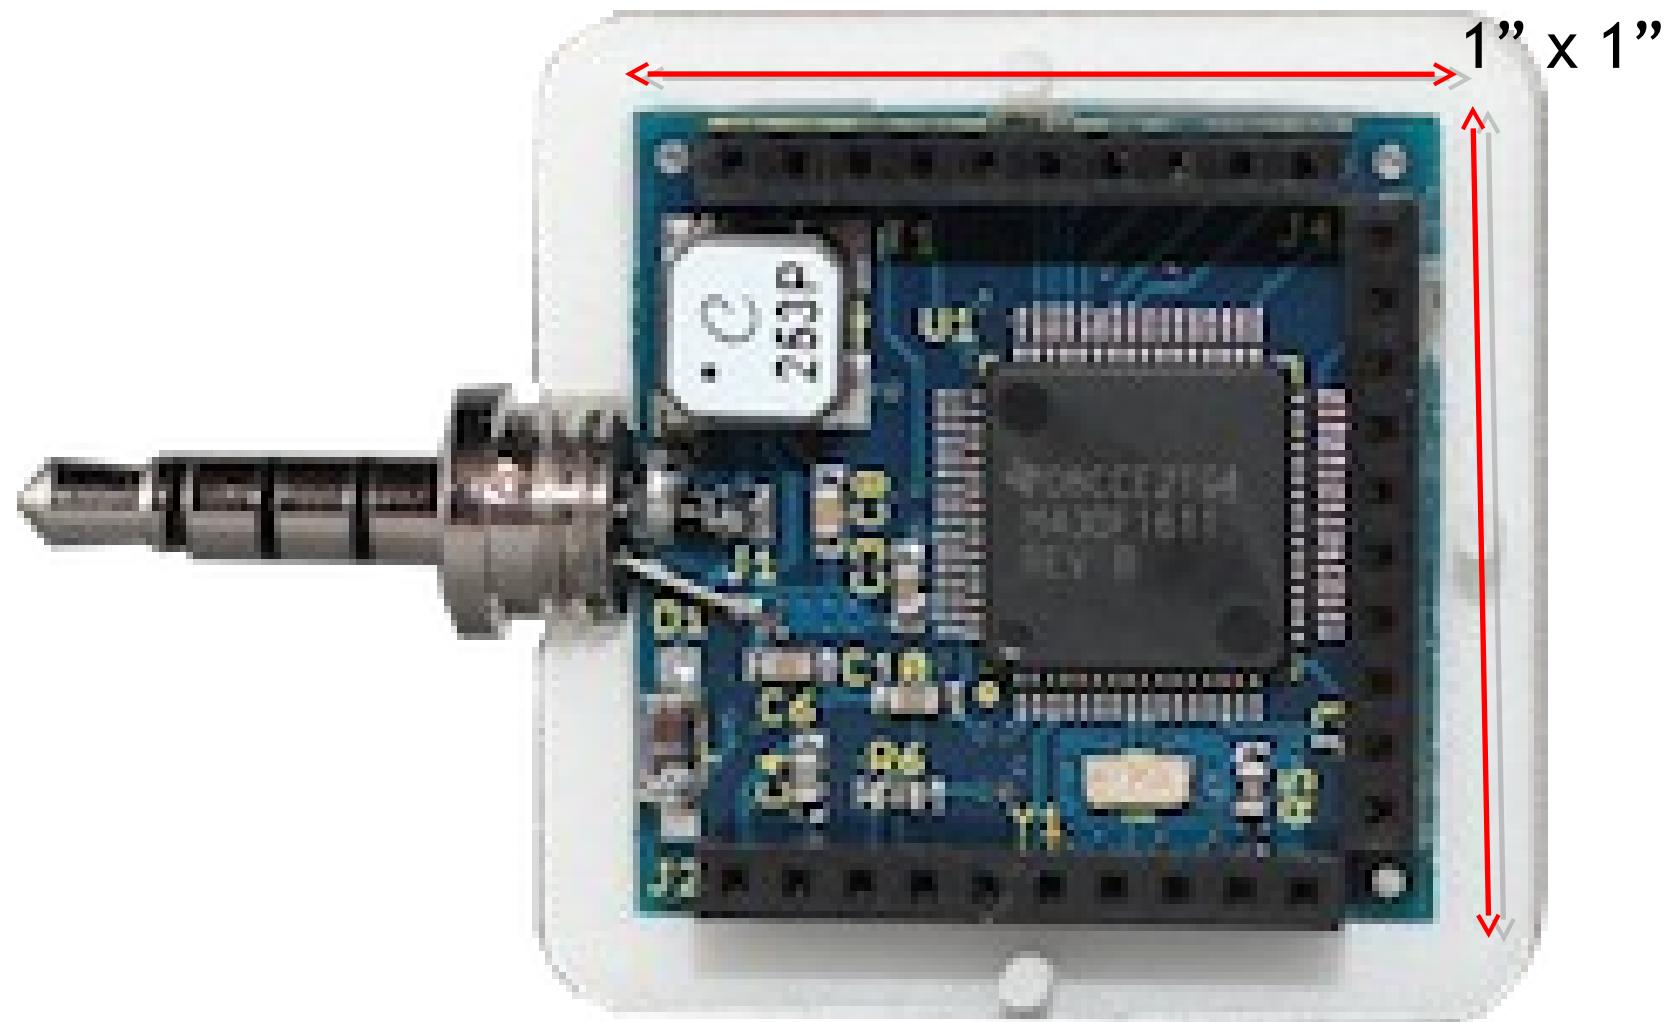


Peripherals  
Peripherals  
Peripherals

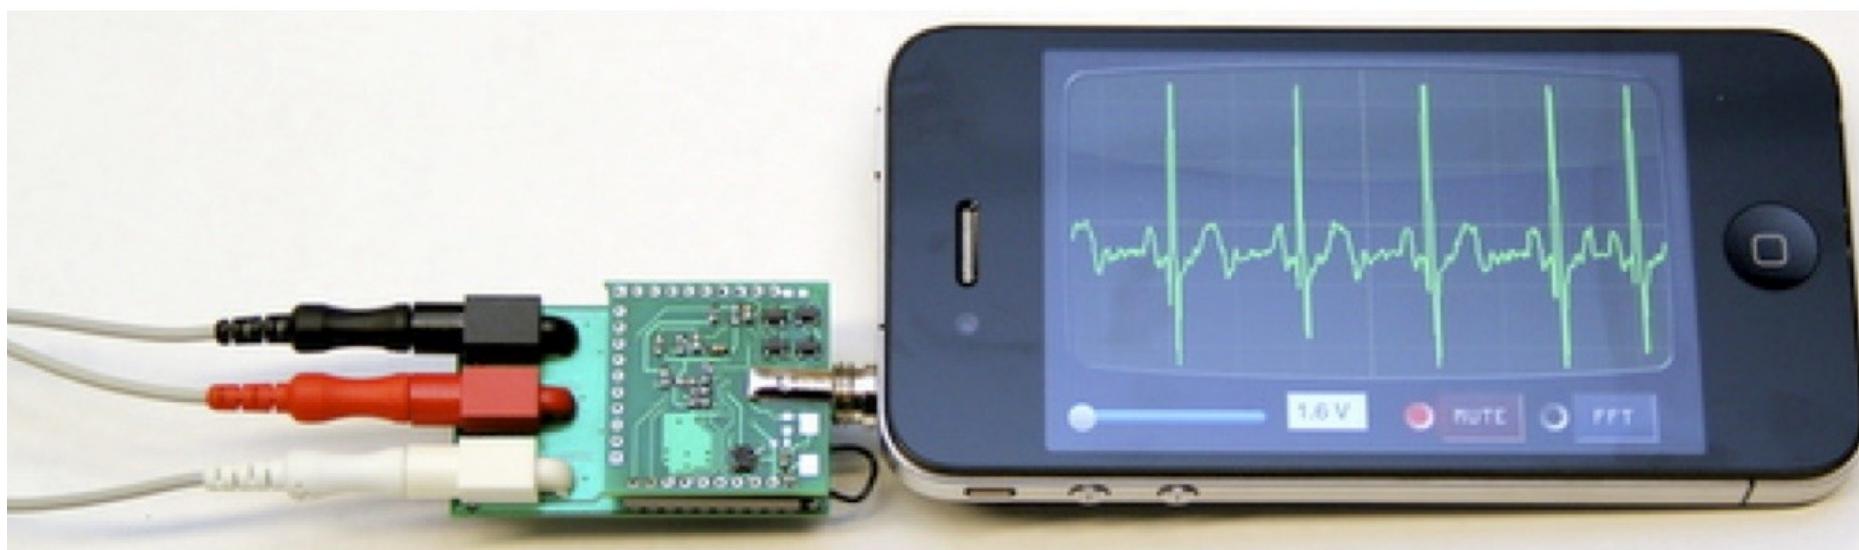


An embedded systems design example:  
Turning the mobile phone into an EKG station

# Integrating power, data, and processing



# Integrating power, data, and processing



# Architecture

In the context of computers,  
what does *architecture* mean?

# Architecture has many meanings

- Computer Organization (or Microarchitecture)
  - Control and data paths
  - Pipeline design
  - Cache design
  - ...
- System Design (or Platform Architecture)
  - Memory and I/O buses
  - Memory controllers
  - Direct memory access
  - ...
- Instruction Set Architecture (ISA)

What is an  
*Instruction Set Architecture (ISA)?*

# Instruction Set Architecture

*“Instruction set architecture (ISA) is the structure of a computer that a machine language programmer (or a compiler) must understand to write a correct (timing independent) program for that machine”*

IBM introducing 360 in 1964

# An ISA defines the hardware/software interface

- A “contract” between architects and programmers
- Register set
- Instruction set
  - Addressing modes
  - Word size
  - Data formats
  - Operating modes
  - Condition codes
- *Calling conventions*
  - Really not part of the ISA (usually)
  - Rather part of the ABI
  - But the ISA often provides meaningful support.

# ARM Architecture roadmap

 <b>4T</b>  ARM7TDMI ARM922T  Thumb instruction set	 <b>5TE</b>  ARM926EJ-S ARM946E-S ARM966E-S  Improved ARM/Thumb Interworking DSP instructions  <b>Extensions:</b> Jazelle (5TEJ)	 <b>6</b>  ARM1136JF-S ARM1176JZF-S ARM11 MPCore  SIMD Instructions Unaligned data support  <b>Extensions:</b> Thumb-2 (6T2) TrustZone (6Z)  Multicore (6K)	 <b>7</b>  Cortex-A8/R4/M3/M1 Thumb-2  <b>Extensions:</b> v7A (applications) – NEON v7R (real time) – HW Divide V7M (microcontroller) – HW Divide and Thumb-2 only
--	--	---	--

# ARM Cortex-M3 ISA

## Instruction Set

ADD Rd, Rn, <op2>

Branching  
Data processing  
Load/Store  
Exceptions  
Miscellaneous

## Register Set

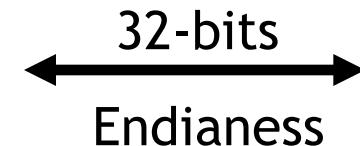
R0
R1
R2
R3
R4
R5
R6
R7
R8
R9
R10
R11
R12
R13 (SP)
R14 (LR)
R15 (PC)
xPSR



Endianess

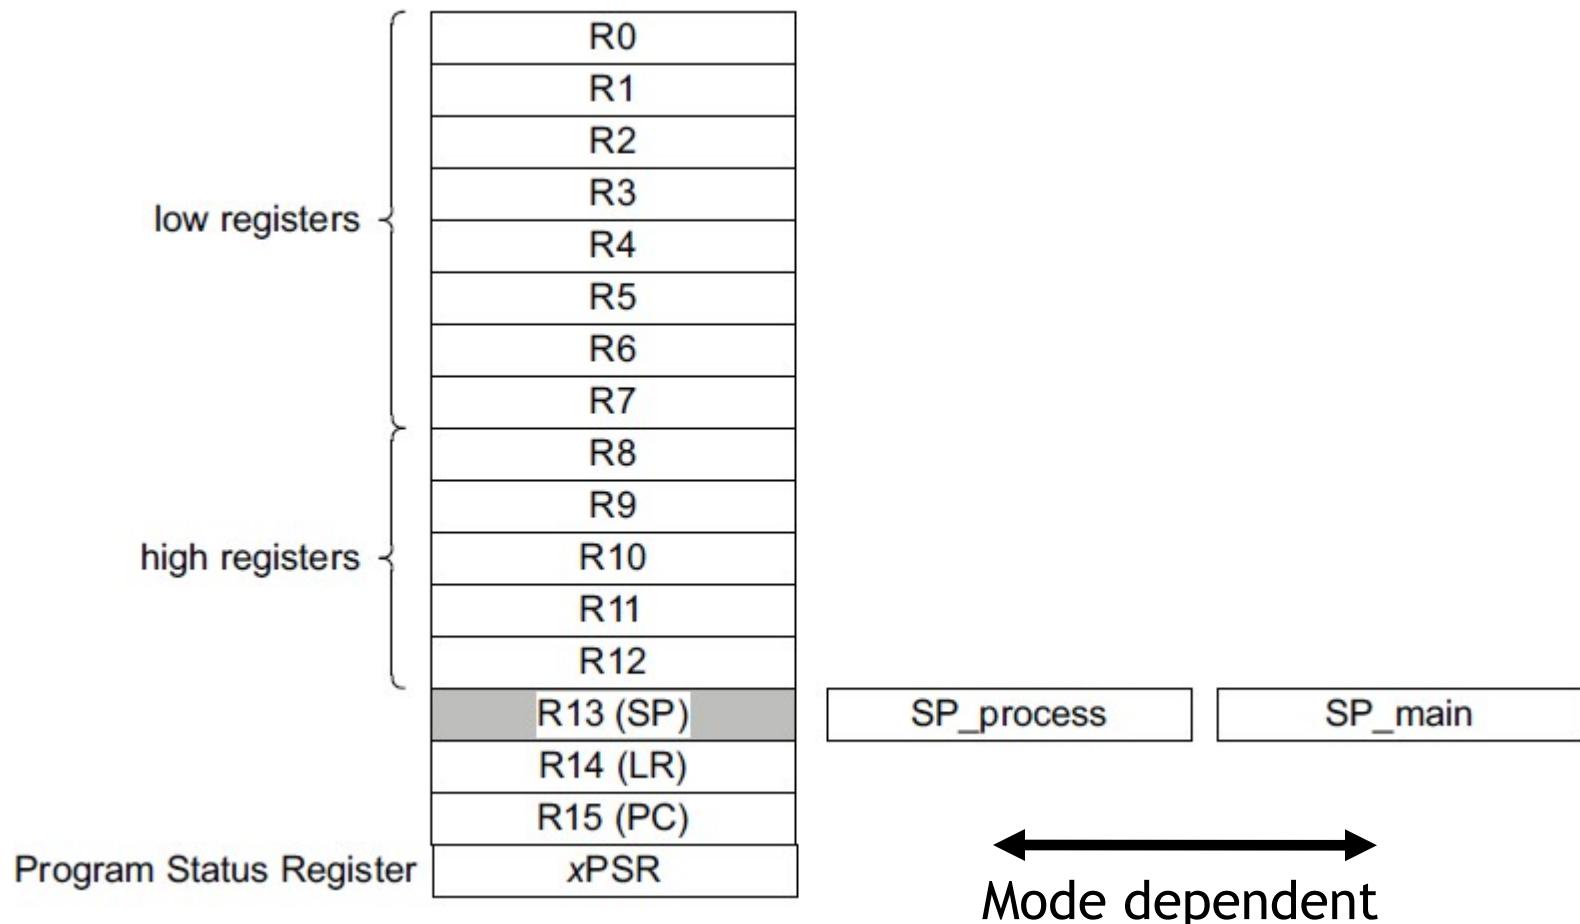
## Address Space

System	0xFFFFFFFF
Private peripheral bus - External	0xE0100000
Private peripheral bus - Internal	0xE0040000
External device 1.0GB	0xE0000000
External RAM 1.0GB	0xA0000000
Peripheral 0.5GB	0x60000000
SRAM 0.5GB	0x40000000
Code 0.5GB	0x20000000
	0x00000000

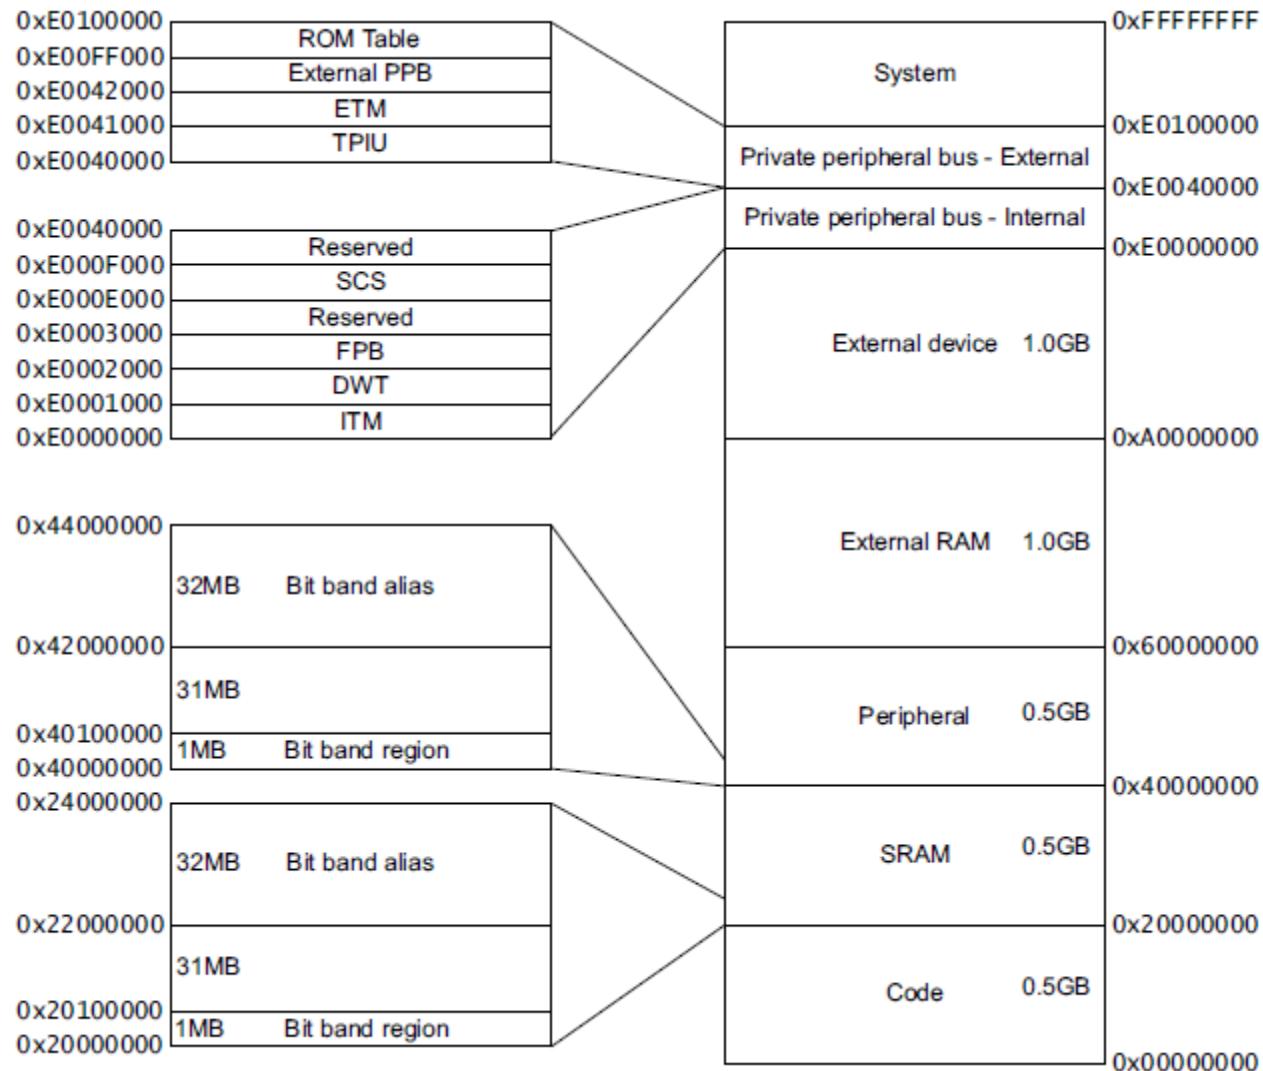


Endianess

# Registers



# Address Space



# Instruction Encoding

## ADD immediate

### Encoding T1 All versions of the Thumb ISA.

ADDS <Rd>,<Rn>,#<imm3>

ADD<c> <Rd>,<Rn>,#<imm3>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	1	1	1	0	imm3	Rn	Rd						

### Encoding T2 All versions of the Thumb ISA.

ADDS <Rdn>,#<imm8>

ADD<c> <Rdn>,#<imm8>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	1	1	0	Rdn		imm8								

### Encoding T3 ARMv7-M

ADD{S}<c>.W <Rd>,<Rn>,#<const>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
1	1	1	1	0	i	0	1	0	0	0	S	Rn	0	imm3	Rd		imm8															

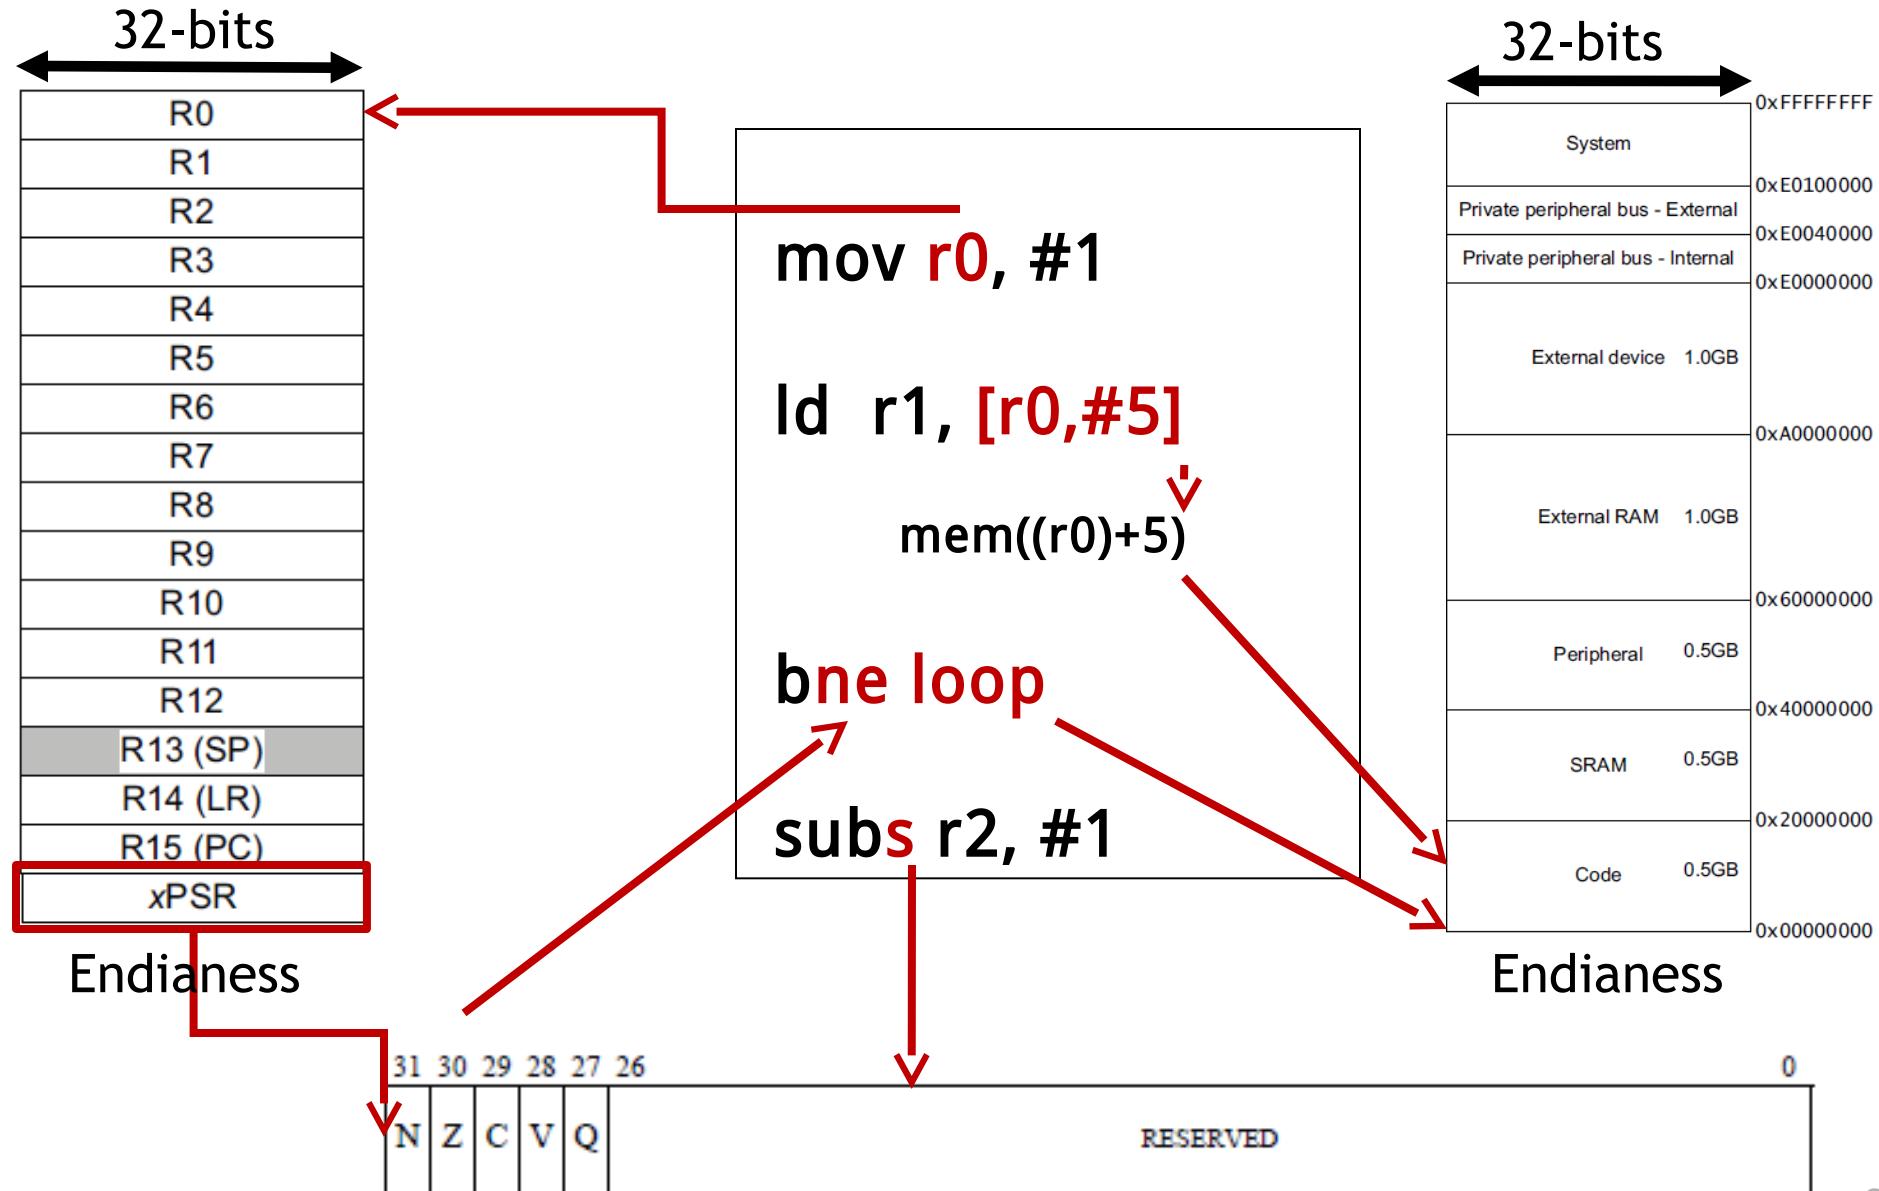
### Encoding T4 ARMv7-M

ADDW<c> <Rd>,<Rn>,#<imm12>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0		
1	1	1	1	0	i	1	0	0	0	0	0	Rn	0	imm3	Rd		imm8																

# Major elements of an Instruction Set Architecture

(registers, memory, word size, endianess, conditions, instructions, addressing modes)



# Instruction classes

- Branching
- Data processing
- Load/store
- Exceptions
- Miscellaneous

# Addressing Modes

- Offset Addressing
  - Offset is added or subtracted from base register
  - Result used as effective address for memory access
  - [ $<Rn>$ ,  $<\text{offset}>$ ]
- Pre-indexed Addressing
  - Offset is applied to base register
  - Result used as effective address for memory access
  - Result written back into base register
  - [ $<Rn>$ ,  $<\text{offset}>$ ]!
- Post-indexed Addressing
  - The address from the base register is used as the EA
  - The offset is applied to the base and then written back
  - [ $<Rn>$ ],  $<\text{offset}>$

## <offset> options

- An immediate constant
  - #10
- An index register
  - <Rm>
- A shifted index register
  - <Rm>, LSL #<shift>

# Updating the Application Program Status Register (aka condition codes or APRS)

- `sub r0, r1`
  - $r0 \leftarrow r0 - r1$
  - APSR remain unchanged
- `subs r0, r1`
  - $r0 \leftarrow r0 - r1$
  - APSR N or Z bits could change
- `add r0, r1`
  - $r0 \leftarrow r0 + r1$
  - APSR remain unchanged
- `adds r0, r1`
  - $r0 \leftarrow r0 + r1$
  - APSR C or V bits could change

# What does some real assembly look like?

0000017c <main>:

17c:	b580	push	{r7, lr}
17e:	b084	sub	sp, #16
180:	af00	add	r7, sp, #0
182:	f04f 0328	mov.w	r3, #40 ; 0x28
186:	60bb	str	r3, [r7, #8]
188:	f04f 0300	mov.w	r3, #0
18c:	60fb	str	r3, [r7, #12]
18e:	f04f 0300	mov.w	r3, #0
192:	603b	str	r3, [r7, #0]
194:	e010	b.n	1b8 <main+0x3c>
196:	6838	ldr	r0, [r7, #0]
198:	f7ff ffb8	bl	10c <factorial>
19c:	4603	mov	r3, r0
19e:	607b	str	r3, [r7, #4]
1a0:	f646 5010	movw	r0, #27920 ; 0x6d10
1a4:	f2c0 0000	movt	r0, #0
1a8:	6839	ldr	r1, [r7, #0]
1aa:	687a	ldr	r2, [r7, #4]
1ac:	f000 f840	bl	230 <printf>
1b0:	683b	ldr	r3, [r7, #0]
...			

# The endianess religious war: 284 years and counting!

- Modern version
  - Danny Cohen
  - IEEE Computer, v14, #10
  - Published in 1981
  - Satire on CS religious war
- Historical Inspiration
  - Jonathan Swift
  - *Gullivers Travels*
  - Published in 1726
  - Satire on Henry-VIII's split with the Church
- Little-Endian
  - LSB is at lower address

Memory Offset	Value (LSB)	(MSB)
0x0000	01 02 FF	00
0x0004	78 56 34	12

```
uint8_t a = 1;  
uint8_t b = 2;  
uint16_t c = 255; // 0x00FF  
uint32_t d = 0x12345678;
```
- Big-Endian
  - MSB is at lower address

Memory Offset	Value (LSB)	(MSB)
0x0000	01 02 00	FF
0x0004	12 34 56	78

```
uint8_t a = 1;  
uint8_t b = 2;  
uint16_t c = 255; // 0x00FF  
uint32_t d = 0x12345678;
```



# Instruction encoding

- Instructions are encoded in machine language opcodes
- Sometimes
  - Necessary to hand generate opcodes
  - Necessary to verify assembled code is correct
- How?

Instructions  
**movs r0, #10**  
**movs r1, #0**

<u>Register Value</u>	<u>Memory Value</u>
<u>001 00 000 00001010</u> (msb)	(LSB) (MSB) <u>0a 20 00 21</u>
<u>001 00 001 00000000</u>	

ARMv7 ARM

Encoding T1

MOV<C> <Rd>, #<imm8>

MOV<C> <Rd>, #<imm8>

All versions of the Thumb ISA.

Outside IT block.

Inside IT block.

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	1	0	0	Rd		imm8								

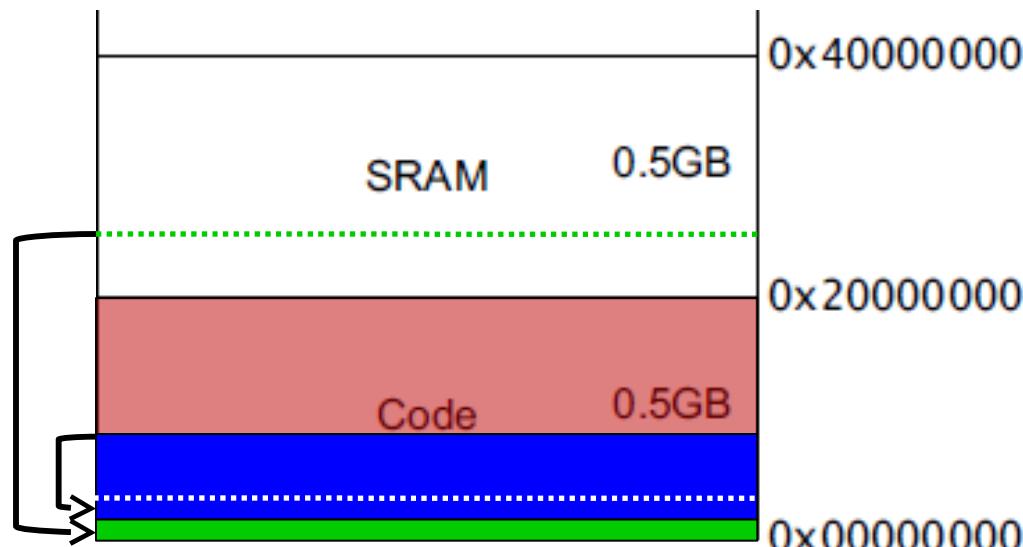
d = UInt(Rd); setflags = !InITBlock(); imm32 = ZeroExtend(imm8, 32); carry = APSR.C;

# What happens after a power-on-reset (POR)?

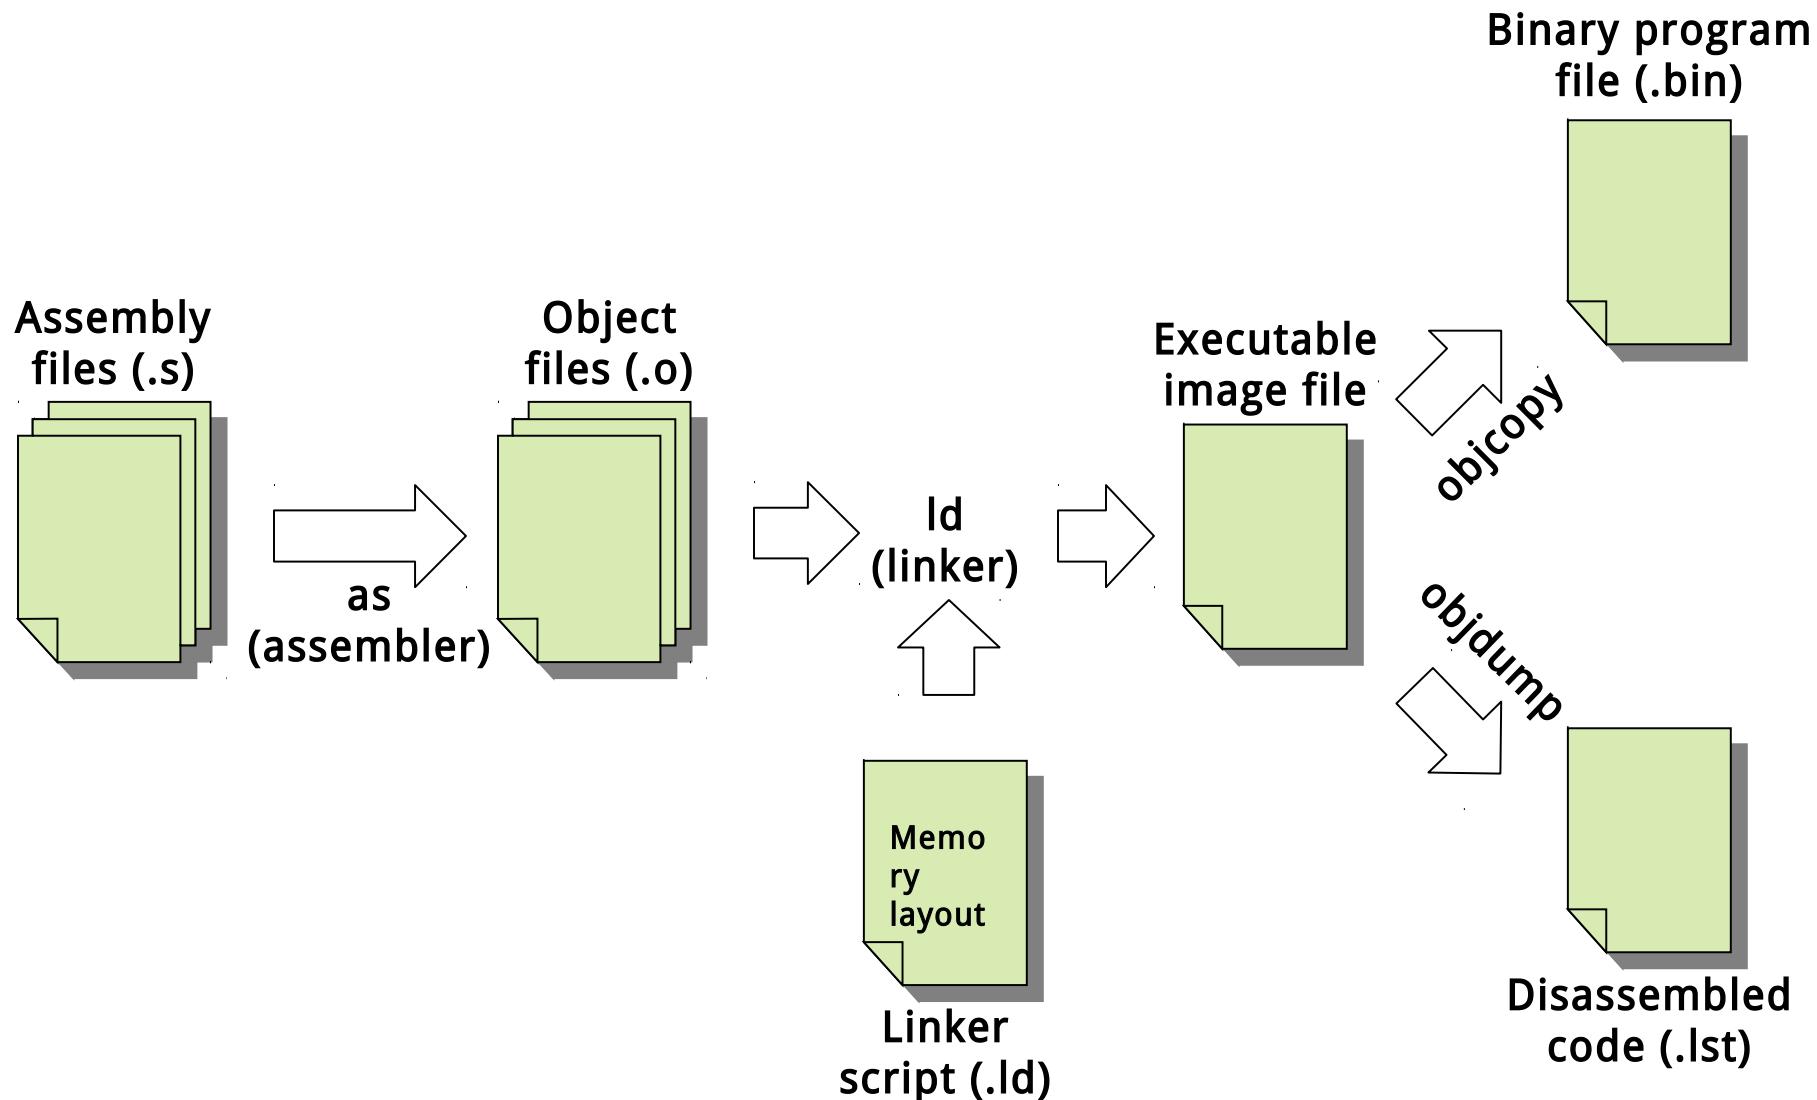
- On the ARM Cortex-M3
- SP and PC are loaded from the code (.text) segment
- **Initial stack pointer**
  - LOC: 0x00000000
  - POR: SP  $\leftarrow$  mem(0x00000000)
- **Interrupt vector table**
  - *Initial* base: 0x00000004
  - Vector table is relocatable
  - Entries: 32-bit values
  - Each entry is an address
  - Entry #1: reset vector
    - LOC: 0x00000004
    - POR: PC  $\leftarrow$  mem(0x00000004)
- **Execution begins**

```
.equ      STACK_TOP, 0x20000800
.text
.syntax unified
.thumb
.global _start
.type    start, %function
```

```
_start:
.start:   .word    STACK_TOP, start
           movs r0, #10
           ...
...
```



# How does an assembly language program get turned into a executable program image?



# What are the real GNU executable names for the ARM?

- Just add the prefix “arm-none-eabi-” prefix
- Assembler (as)
  - arm-none-eabi-as
- Linker (ld)
  - arm-none-eabi-ld
- Object copy (objcopy)
  - arm-none-eabi-objcopy
- Object dump (objdump)
  - arm-none-eabi-objdump
- C Compiler (gcc)
  - arm-none-eabi-gcc
- C++ Compiler (g++)
  - arm-none-eabi-g++

# A simple (hardcoded) Makefile example

```
all:
```

```
    arm-none-eabi-as -mcpu=cortex-m3 -mthumb example1.s -o example1.o  
    arm-none-eabi-ld -Ttext 0x0 -o example1.out example1.o  
    arm-none-eabi-objcopy -Obinary example1.out example1.bin  
    arm-none-eabi-objdump -S example1.out > example1.lst
```

# What information does the disassembled file provide?

all:

```
arm-none-eabi-as -mcpu=cortex-m3 -mthumb example1.s -o example1.o  
arm-none-eabi-ld -Ttext 0x0 -o example1.out example1.o  
arm-none-eabi-objcopy -Obinary example1.out example1.bin  
arm-none-eabi-objdump -S example1.out > example1.lst
```

```
.equ      STACK_TOP, 0x20000800  
.text  
.syntax unified  
.thumb  
.global _start  
.type    start, %function  
  
.word     STACK_TOP, start  
  
.start:  
start:    movs r0, #10  
          movs r1, #0  
  
.loop:  
loop:     adds r1, r0  
          subs r0, #1  
          bne  loop  
  
.deadloop:  
deadloop: b  deadloop  
.end
```

example1.out: file format elf32-littlearm

Disassembly of section .text:

00000000 <\_start>:  
0: 20000800 .word 0x20000800  
4: 00000009 .word 0x00000009

00000008 <start>:  
8: 200a movs r0, #10  
a: 2100 movs r1, #0

0000000c <loop>:  
c: 1809 adds r1, r0  
e: 3801 subs r0, #1  
10: d1fc bne.n c <loop>

00000012 <deadloop>:  
12: e7fe b.n 12  
<deadloop>

# What are the elements of a real assembly program?

```
.equ      STACK_TOP, 0x20000800      /* Equates symbol to value */
.text
.syntax unified                      /* Tells AS to assemble region */
.thumb
.global _start                        /* Means language is ARM UAL */
                                    /* Means ARM ISA is Thumb */
                                    /* .global exposes symbol */
                                    /* _start label is the beginning */
                                    /* ...of the program region */
                                    /* Specifies start is a function */
                                    /* start label is reset handler */

_start:
.word     STACK_TOP, start          /* Inserts word 0x20000800 */
                                    /* Inserts word (start) */

start:
    movs r0, #10                    /* We've seen the rest ... */
    movs r1, #0

loop:
    adds r1, r0
    subs r0, #1
    bne  loop

deadloop:
    b    deadloop

.end
```

# How are assembly files assembled?

- \$ arm-none-eabi-as
  - Useful options
    - -mcpu
    - -mthumb
    - -O

```
$ arm-none-eabi-as -mcpu=cortex-m3 -mthumb example1.s -o example1.o
```

# How can the contents of an object file be read?

- \$ readelf -a example.o
- Shows
  - ELF headers
  - Program headers
  - Section headers
  - Symbol table
  - Files attributes
- Other options
  - -s shows symbols
  - -S shows section headers

# What does an object file contain?

```
$ readelf -S example1.o
```

There are **9 section headers**, starting at offset 0xac:

## Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	00000000	0000000	00000000	00		0	0	0
[ 1]	.text	PROGBITS	00000000	000034	000014	00	AX	0	0	1
[ 2]	.rel.text	REL	00000000	000300	000008	08		7	1	4
[ 3]	.data	PROGBITS	00000000	000048	0000000	00	WA	0	0	1
[ 4]	.bss	NOBITS	00000000	000048	0000000	00	WA	0	0	1
[ 5]	.ARM.attributes	ARM_ATTRIBUTES	00000000	000048	000021	00		0	0	1
[ 6]	.shstrtab	STRTAB	00000000	000069	000040	00		0	0	1
[ 7]	.symtab	SYMTAB	00000000	000214	0000c0	10		8	11	4
[ 8]	.strtab	STRTAB	00000000	0002d4	00002c	00		0	0	1

## Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings)

I (info), L (link order), G (group), x (unknown)

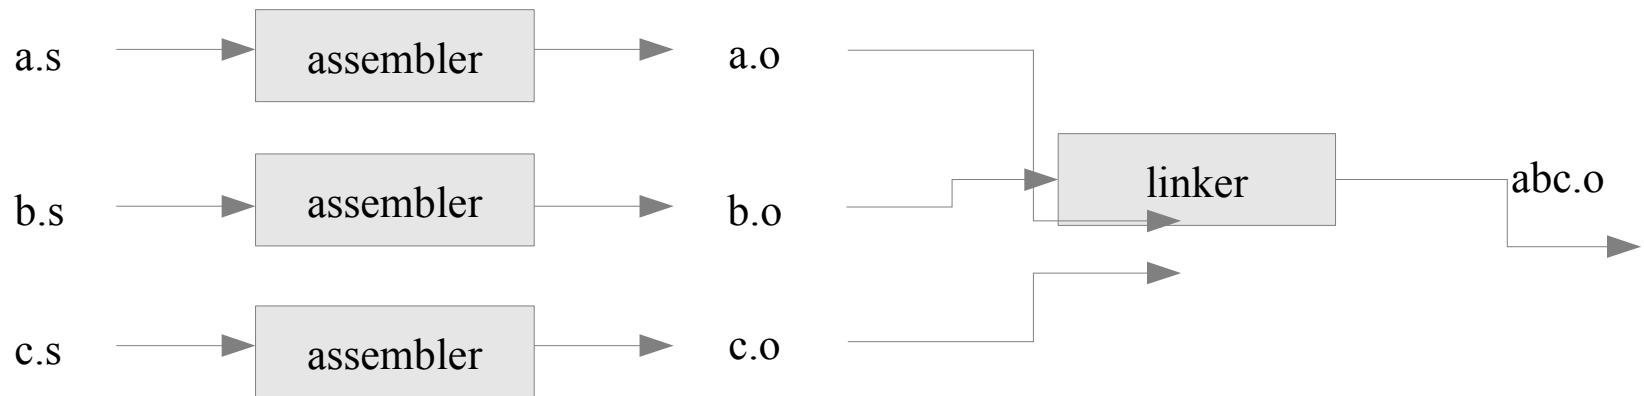
O (extra OS processing required) o (OS specific), p (processor specific)

# How are object files linked?

- \$ arm-none-eabi-ld
  - Useful options
    - -Ttext
    - -Tbss
    - -O

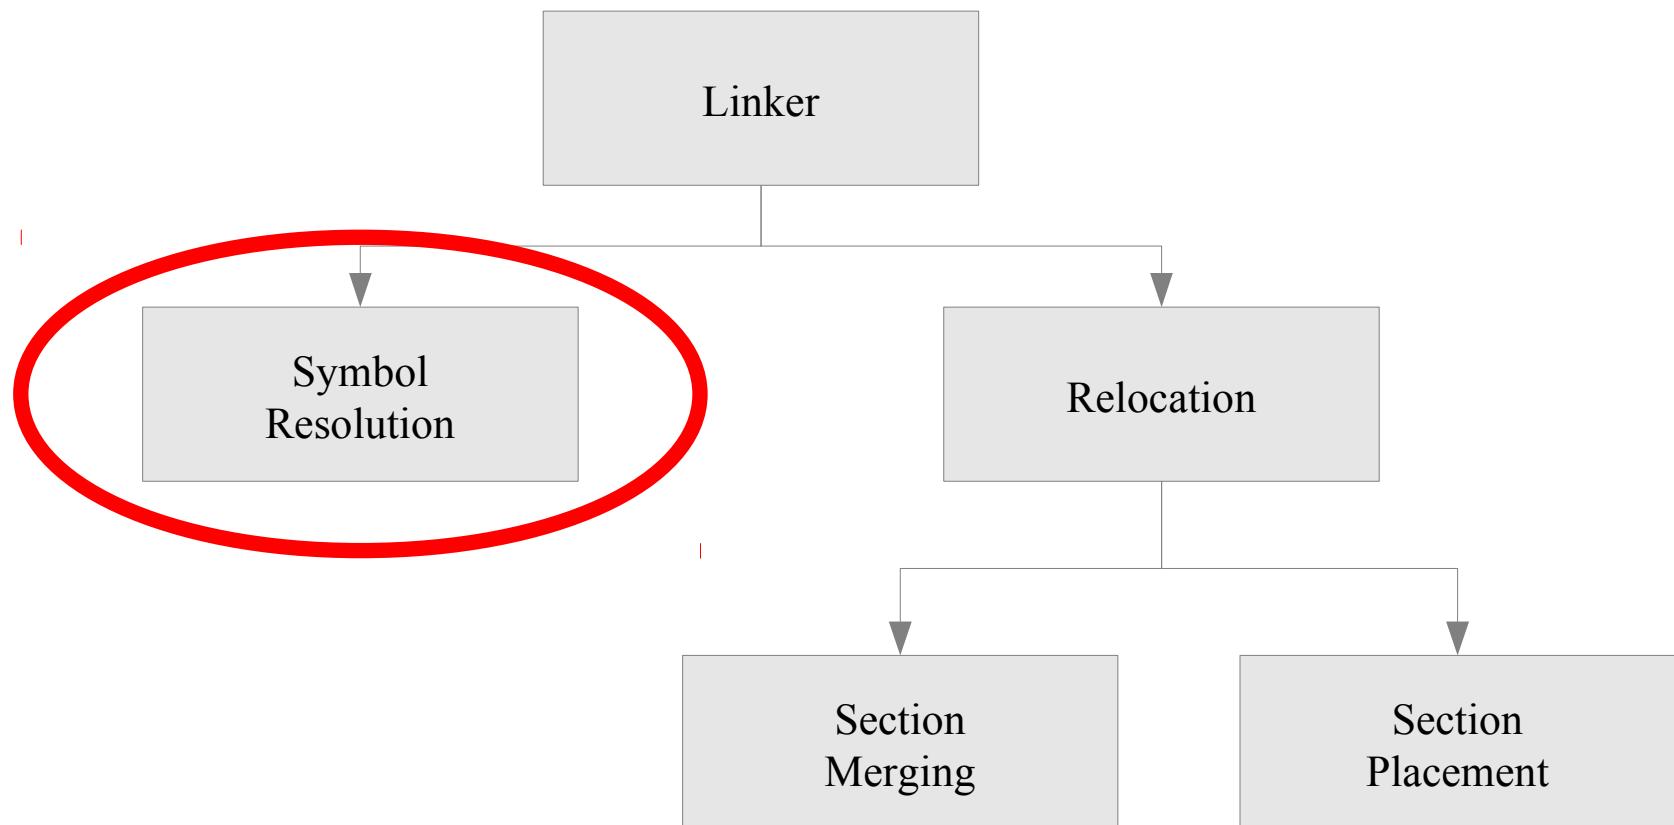
```
$ arm-none-eabi-ld -Ttext 0x0 -Tbss 0x20000000 -o example1.out example1.o
```

# Linker



- In multi-file programs - combines multiple object files to form executable

# Linker

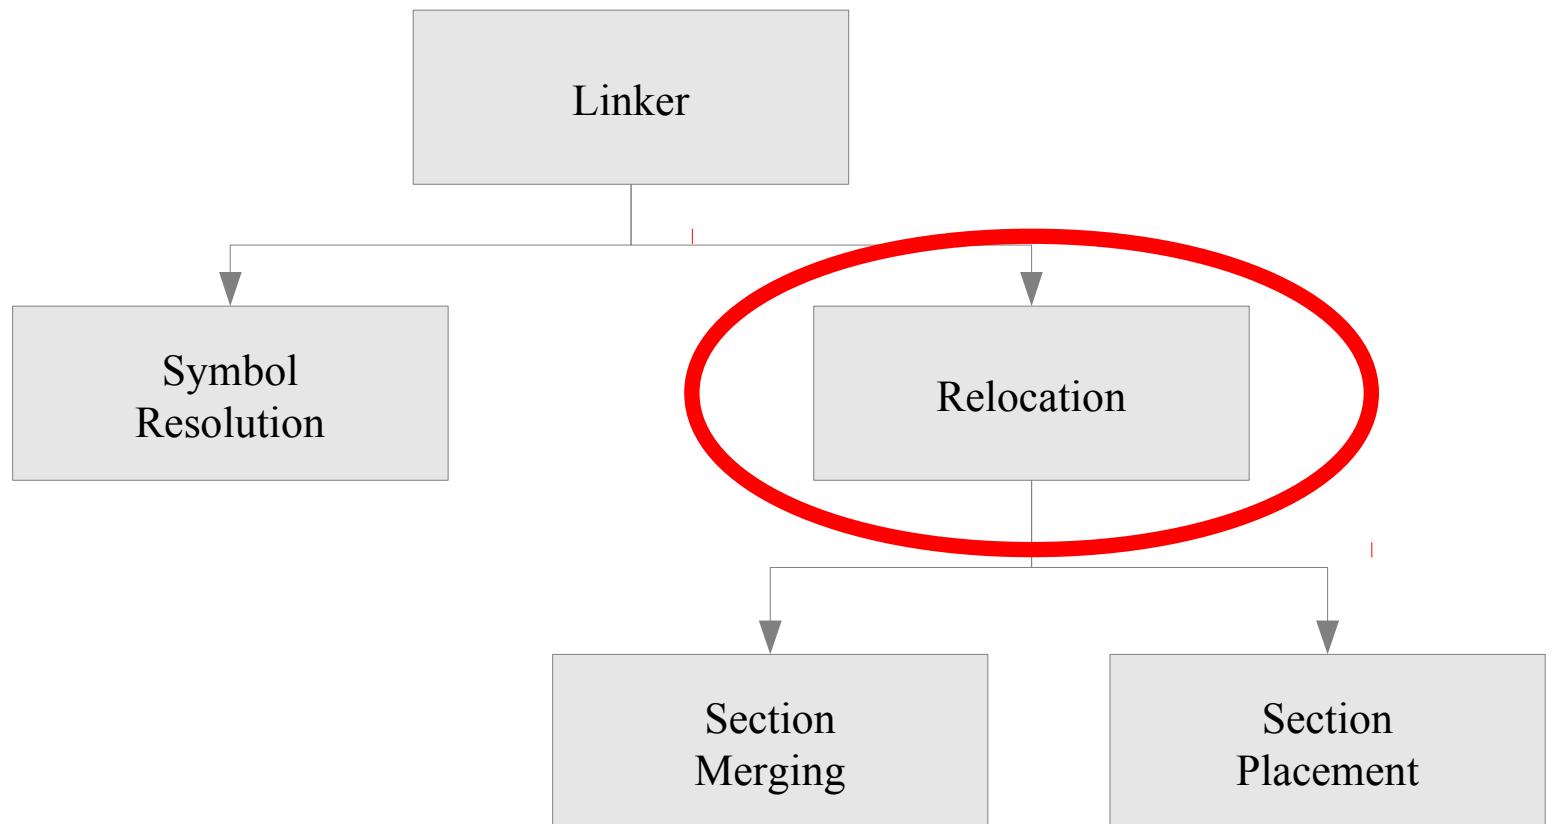


# Symbol Resolution

- Functions are defined in one file
- Referenced in another file
- References are marked unresolved by the compiler
- Linker patches the references



# Linker



# Relocation

- Code generated assuming it starts from address X
- Code should start from address Y
- Change addresses assigned to labels
- Patch label references



# Sections

- Placing related bytes at a particular location.
- Example:
  - instructions in Flash
  - data in RAM
- Related bytes are grouped together using sections
- Placement of sections can be specified

# Sections

- Most programs have atleast two sections, .text and .data
- Data or instructions can be placed in a section using directives
- Directives
  - .text
  - .data
  - .section

# Sections

```
.data  
arr: .word 10, 20, 30, 40, 50  
len: .word 5  
.text  
start: mov r1, #10  
       mov r2, #20  
.data  
result: .skip 4  
.text  
       add r3, r2, r1  
       sub r3, r2, r1
```

## .data section

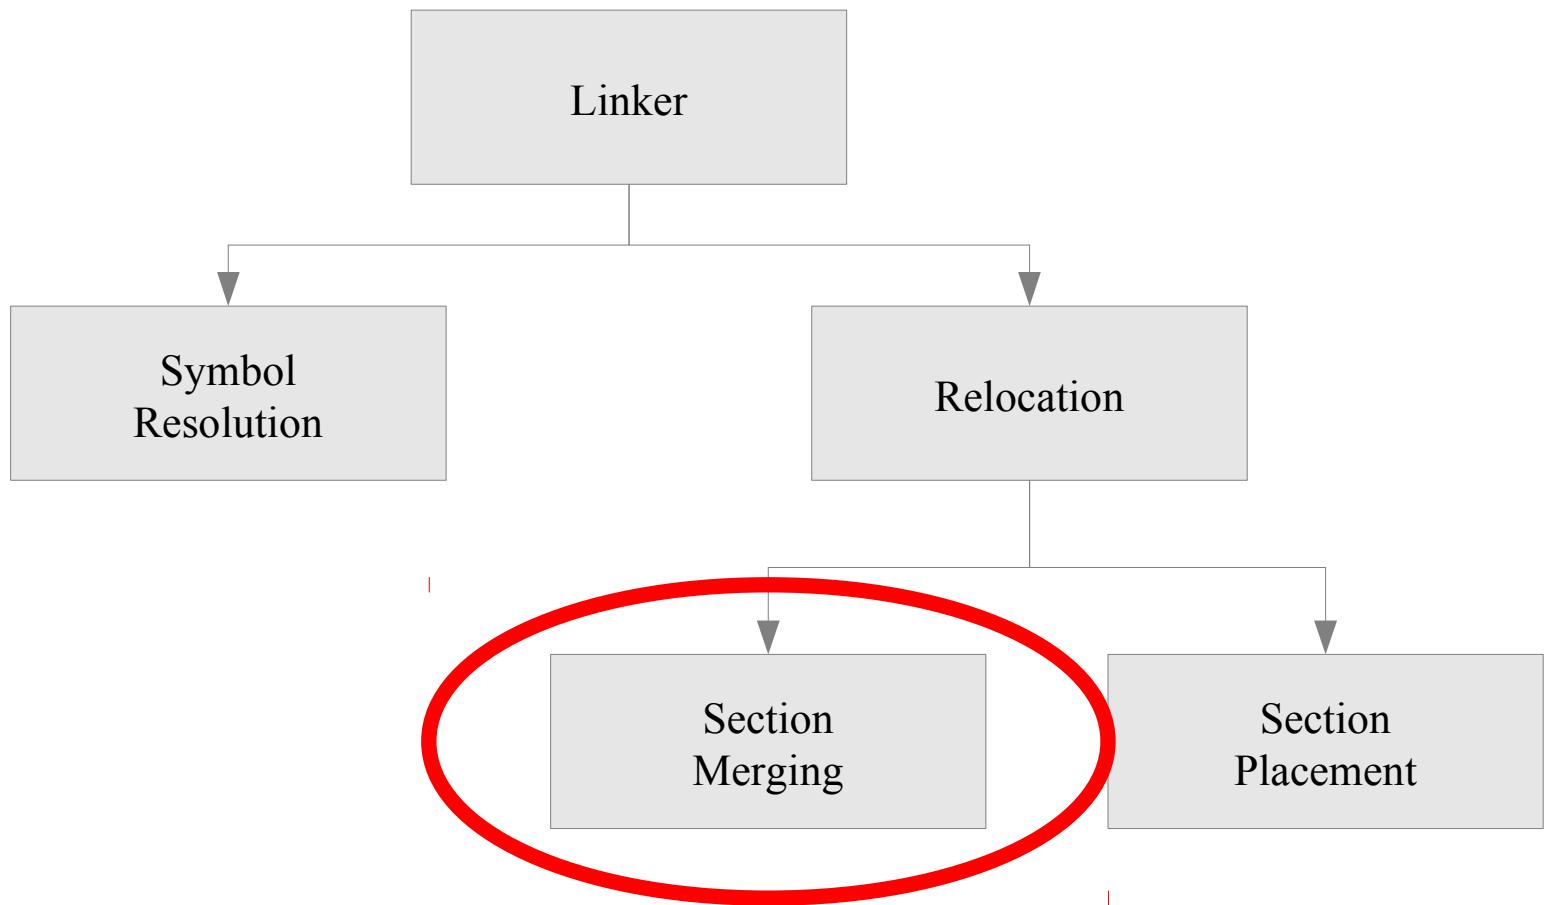
```
0000_0000 arr: .word 10, 20, 30, 40, 50  
0000_0014 len: .word 5  
0000_0018 result: .skip 4
```

## .text section

```
0000_0000 start: mov r1, #10  
0000_0004           mov r2, #20  
0000_0008           add r3, r2, r1  
0000_000C           sub r3, r2, r1
```

- Source - sections can be interleaved
- Bytes of a section - contiguous addresses

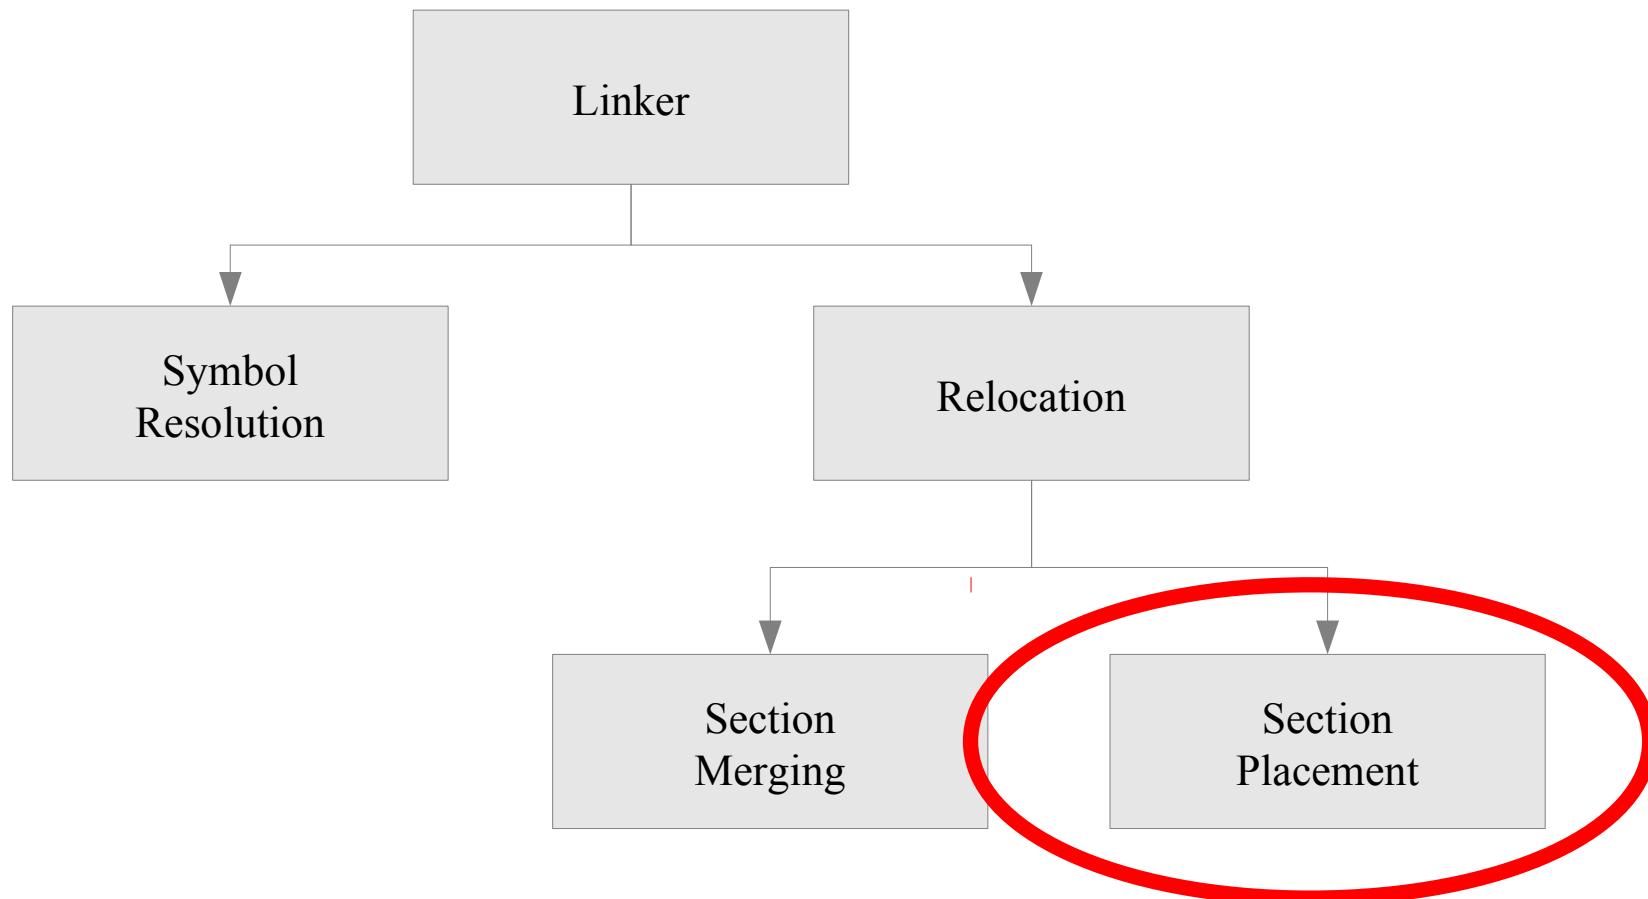
# Linker



## Section Merging

- Linker merges sections in the input files into sections in the output file
- Default merging - sections of same name
- Symbols get new addresses, and references are patched
- Section merging can be controlled by linker script files

# Linker



# Section Placement

- Bytes in each section is given addresses starting from 0x0
- Labels get addresses relative to the start of section
- Linker places section at a particular address
- Labels get new address, label references are patched

a.s (.text)

```
strcpy: ldrb r0, [r1], #1  
       strb r0, [r2], #1  
       cmp  r0, 0  
       bne  strcpy  
       mov   pc, lr
```

Assembler

```
0000_0000 strcpy: ldrb r0, [r1], #1  
0000_0004          strb r0, [r2], #1  
0000_0008          cmp  r0, 0  
0000_000C          bne  strcpy  
0000_0010          mov   pc, lr
```

b.s (.text)

```
strlen: ldrb r0, [r1], #1  
        add   r2, #1  
        cmp   r0, 0  
        bne  strlen  
        mov   pc, lr
```

Assembler

```
0000_0000 strlen: ldrb r0, [r1], #1  
0000_0004          add   r2, #1  
0000_0008          cmp   r0, 0  
0000_000C          bne  strcpy  
0000_0010          mov   pc, lr
```

```
0000_0000 strcpy: ldrb r0, [r1], #1  
0000_0004          strb r0, [r2], #1  
0000_0008          cmp r0, 0  
0000_000C          bne strcpy  
0000_0010          mov pc, lr
```

```
0000_0000 strlen: ldrb r0, [r1], #1  
0000_0004          add r2, #1  
0000_0008          cmp r0, 0  
0000_000C          bne strlen  
0000_0010          mov
```

Merging .text sections from two files

```
0000_0000 strcpy: ldrb r0, [r1], #1  
0000_0004          strb r0, [r2], #1  
0000_0008          cmp r0, 0  
0000_000C          bne strcpy  
0000_0010          mov pc, lr  
0000_0014          strlen: ldrb r0, [r1], #1  
0000_0018          add r2, #1  
0000_001C          cmp r0, 0  
0000_0020          bne strlen  
0000_0024          mov pc, lr
```

New address  
after merge

Patched

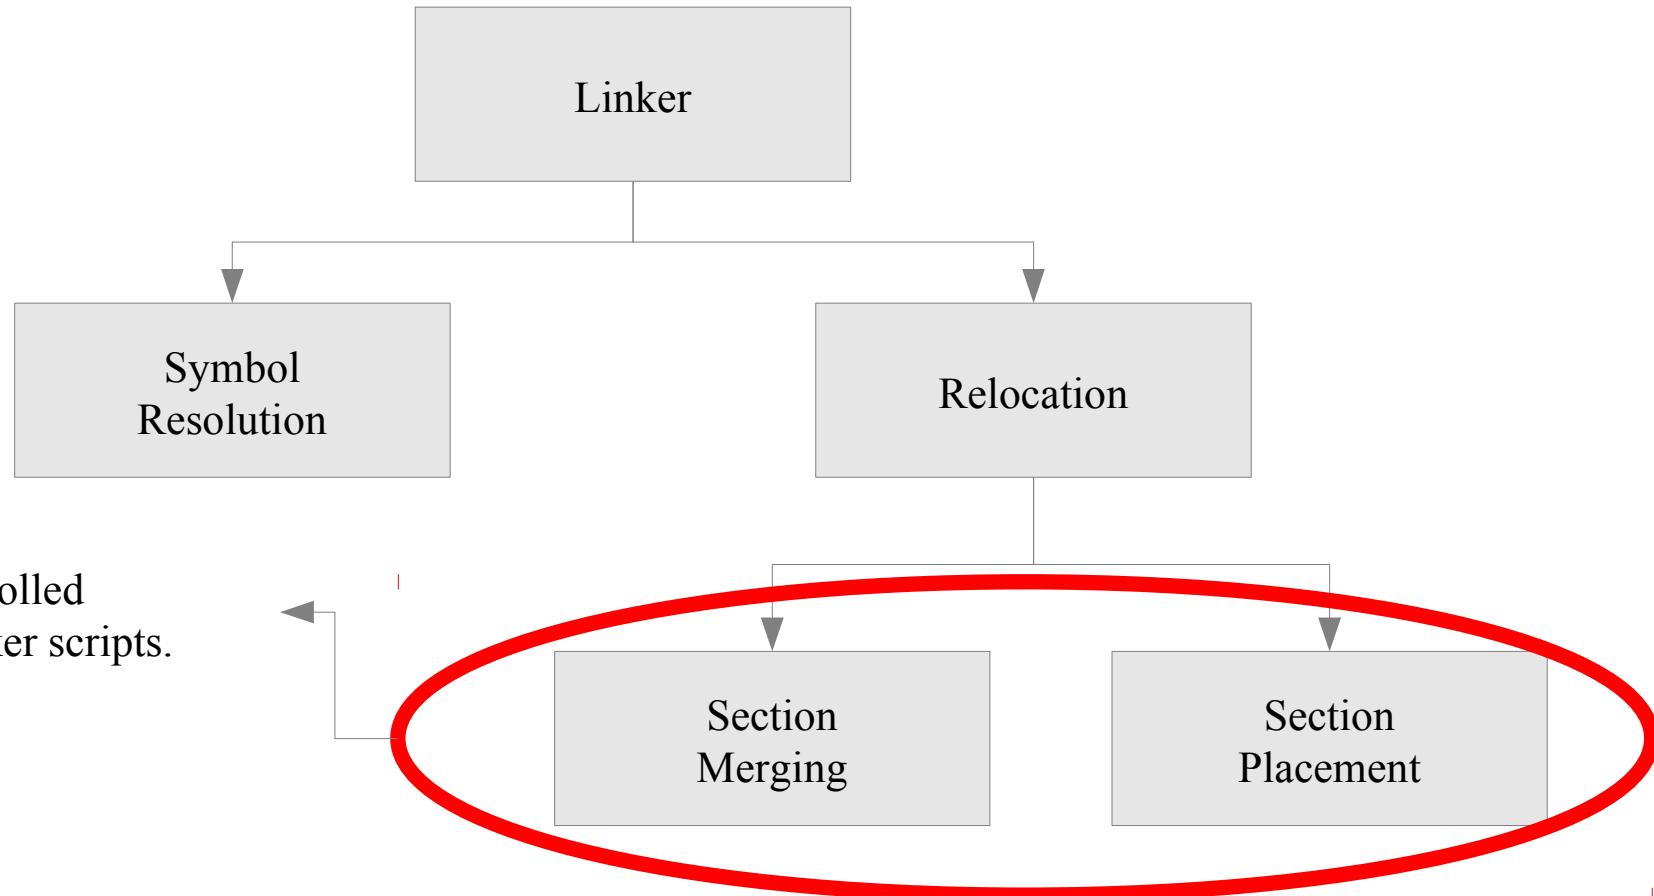
```
0000_0000 strcpy: ldrb r0, [r1], #1  
0000_0004          strb r0, [r2], #1  
0000_0008          cmp  r0, 0  
0000_000C          bne   strcpy  
0000_0010          mov   pc, lr  
0000_0014 strlen: ldrb r0, [r1], #1  
0000_0018          add   r2, #1  
0000_001C          cmp  r0, 0  
0000_0020          bne   strlen  
0000_0024          mov   pc, lr
```

Placing .text section at 0x2000\_0000

```
2000_0000 strcpy: ldrb r0, [r1], #1  
2000_0004          strb r0, [r2], #1  
2000_0008          cmp  r0, 0  
2000_000C          bne   strcpy  
2000_0010          mov   pc, lr  
2000_0014 strlen: ldrb r0, [r1], #1  
2000_0018          add   r2, #1  
2000_001C          cmp  r0, 0  
2000_0020          bne   strlen  
2000_0024          mov   pc, lr
```

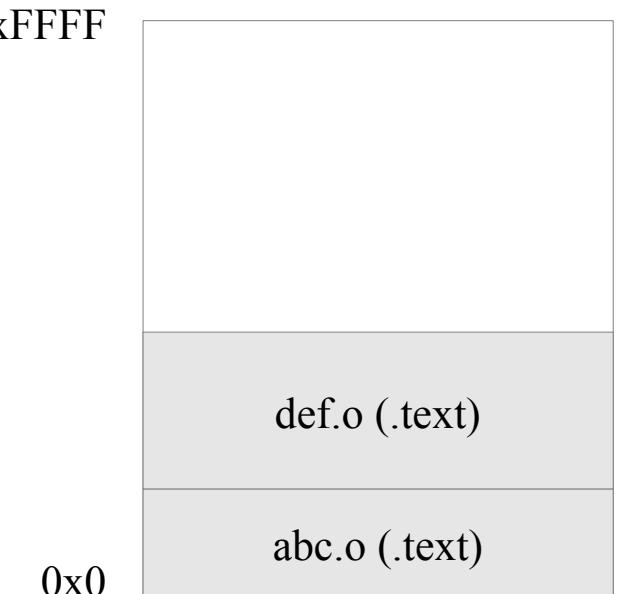
Patched

# Linker Script



# Simple Linker Script

```
MEMORY {  
    FLASH (rx) : ORIGIN = 0x00000000, LENGTH = 0x10000  
    SRAM (rwx) : ORIGIN = 0x20000000, LENGTH = 0x2000  
}  
  
SECTIONS {  
    .text : {  
        abc.o (.text);  
        def.o (.text);  
    } > FLASH  
}
```



# Simple Linker Script

```
MEMORY {  
    FLASH (rx) : ORIGIN = 0x00000000, LENGTH = 0x10000  
    SRAM (rwx) : ORIGIN = 0x20000000, LENGTH = 0x2000  
}  
  
SECTIONS {  
    .text : {  
        abc.o (.text);  
        def.o (.text);  
    } > FLASH  
}
```

→ Section Merging

# Simple Linker Script

```
MEMORY {  
    FLASH (rx) : ORIGIN = 0x00000000, LENGTH = 0x10000  
    SRAM (rwx) : ORIGIN = 0x20000000, LENGTH = 0x2000  
}  
  
SECTIONS {  
    .text : {  
        abc.o (.text);  
        def.o (.text);  
    } > FLASH  
}
```



Section Placement

# Making it Generic

```
MEMORY {  
    FLASH (rx) : ORIGIN = 0x00000000, LENGTH = 0x10000  
    SRAM (rwx) : ORIGIN = 0x20000000, LENGTH = 0x2000  
}  
  
SECTIONS {  
    .text : {  
        * (.text);  
    } > FLASH  
}
```

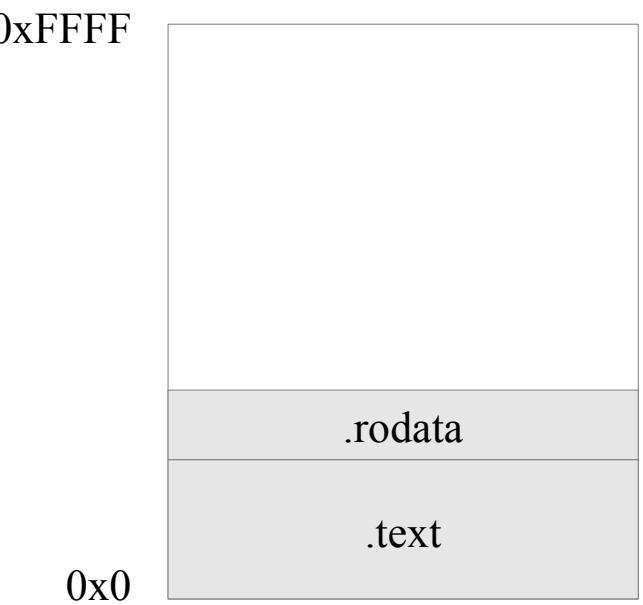
→ Wildcards to represent .text  
form all input files

# Multiple Sections

```
MEMORY {  
    FLASH (rx) : ORIGIN = 0x00000000, LENGTH = 0x10000  
    SRAM (rwx) : ORIGIN = 0x20000000, LENGTH = 0x2000  
}
```

```
SECTIONS {  
    .text : {  
        * (.text);  
    } > FLASH  
  
    .rodata : {  
        * (.rodata);  
    } > FLASH  
}
```

Dealing with  
multiple sections



# RAM is Volatile!

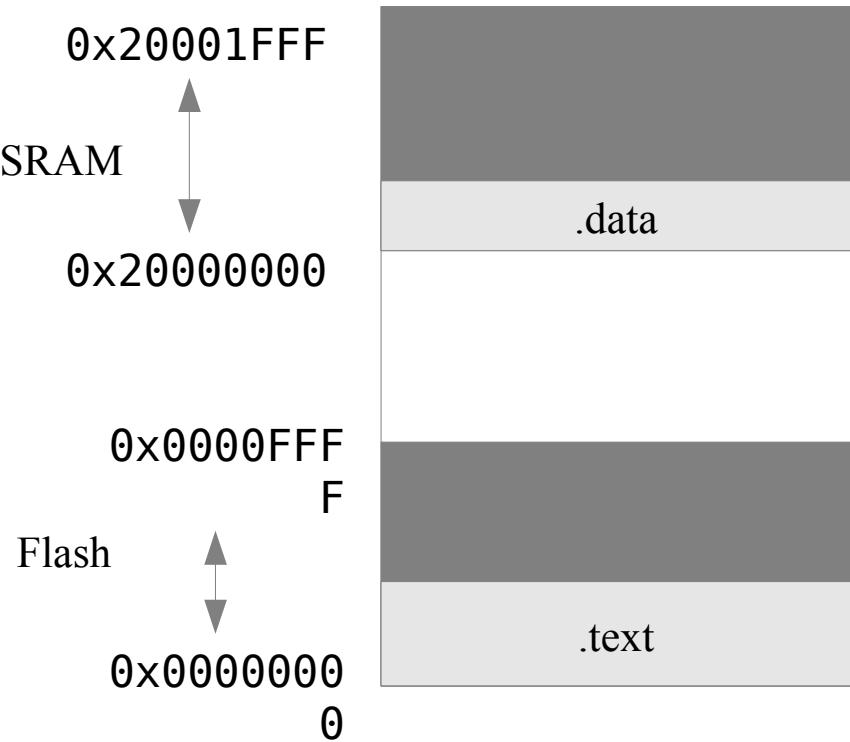
- RAM is volatile
- Data cannot be made available in RAM at power-up
- All code and data should be in Flash at power-up
- Startup code - copies data from Flash to RAM

# RAM is Volatile!

- .data section should be present in Flash at power-up
- Section has two addresses
  - load address (aka LMA)
  - run-time address (aka VMA)
- So far only run-time address - actual address assigned to labels
- Load address defaults to run-time address

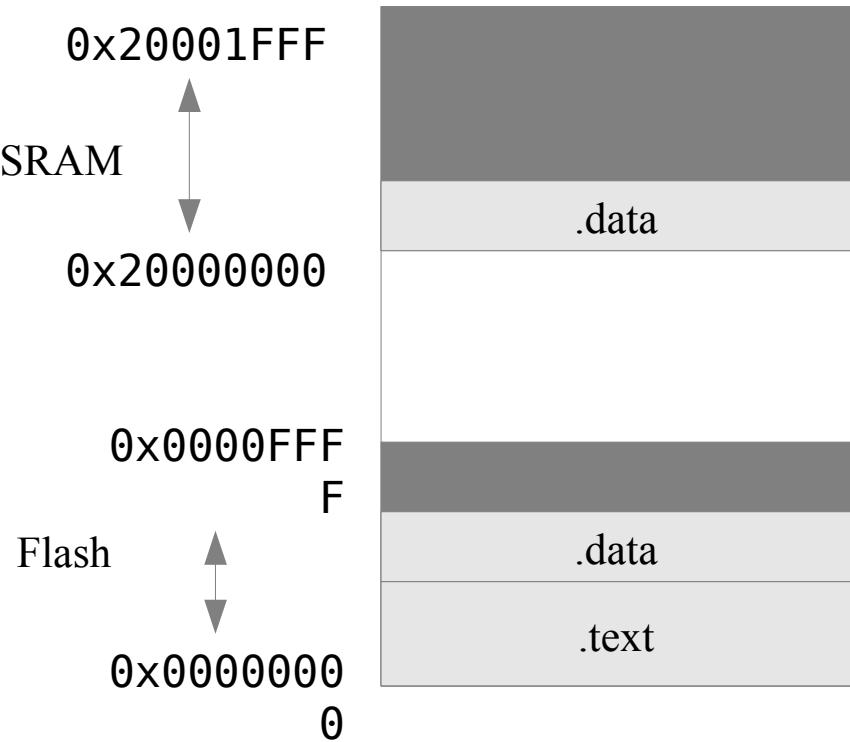
# Linker Script Revisited

```
MEMORY {  
    FLASH (rx) : ORIGIN = 0x00000000, LENGTH = 0x10000  
    SRAM (rwx) : ORIGIN = 0x20000000, LENGTH = 0x2000  
}  
  
SECTIONS {  
    .text : {  
        * (.text);  
    } > FLASH  
  
    .data : {  
        * (.data);  
    } > SRAM  
}
```



# Linker Script Revisited

```
MEMORY {  
    FLASH (rx) : ORIGIN = 0x00000000, LENGTH = 0x10000  
    SRAM (rwx) : ORIGIN = 0x20000000, LENGTH = 0x2000  
}  
  
SECTIONS {  
    .text : {  
        * (.text);  
    } > FLASH  
  
    .data : {  
        * (.data);  
    } > SRAM AT> FLASH  
}
```



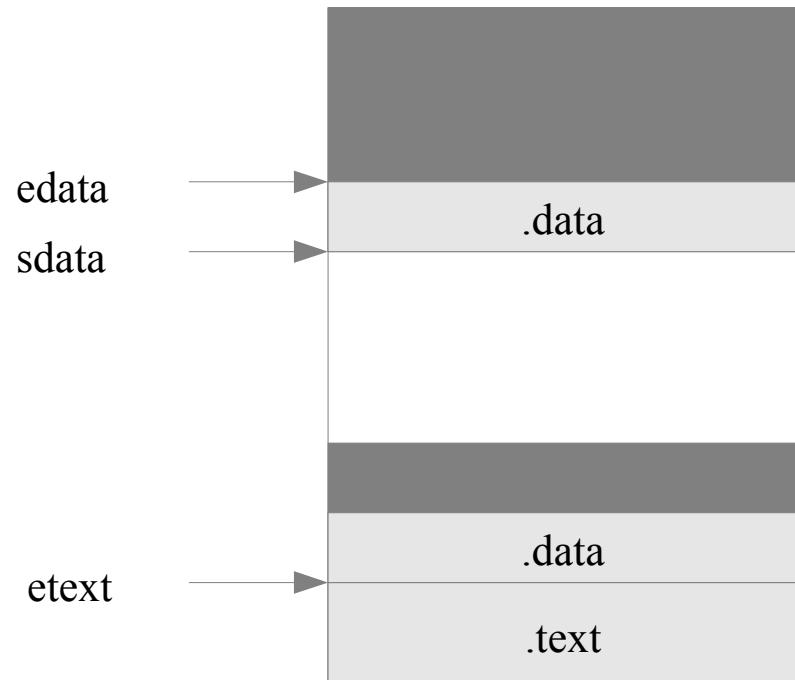
# Linker Script Revisited

```
MEMORY {  
    FLASH (rx) : ORIGIN = 0x00000000, LENGTH = 0x10000  
    SRAM (rwx) : ORIGIN = 0x20000000, LENGTH = 0x2000
```

```
}
```

```
SECTIONS {  
    .text : {  
        * (.text);  
        etext = .;  
    } > FLASH
```

```
.data : {  
    sdata = .;  
    * (.data);  
    edata = .;  
} > SRAM AT> FLASH
```



# Data in RAM

- Copy .data from Flash to RAM

```
start:          ldr r0, =sdata           @ Load the address of sdata
                ldr r1, =edata            @ Load the address of edata
                ldr r2, =etext            @ Load the address of etext

copy:          ldrb r3, [r2]           @ Load the value from Flash
                strb r3, [r0]            @ Store the value in RAM

                add r2, r2, #1          @ Increment Flash pointer
                add r0, r0, #1          @ Increment RAM pointer

                cmp r0, r1             @ Check if end of data
                bne copy               @ Branch if not end of data
```

Lab-1

**`libraries/CMSIS/CM3/DeviceSupport/ST/STM32F10x/startup/gcc_ride7/startup_stm32f10x_md.s`**

# What are the contents of typical linker script?

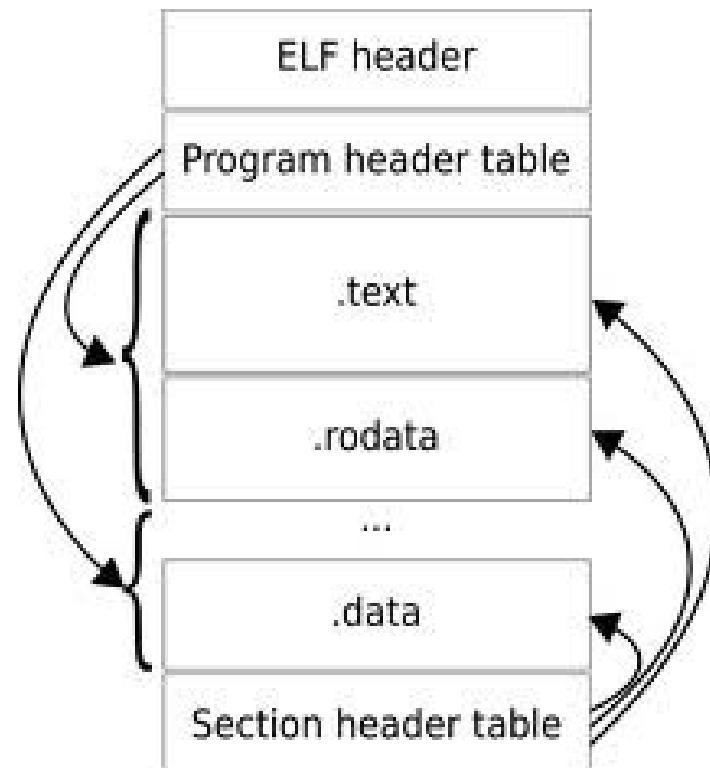
```
OUTPUT_FORMAT("elf32-littlearm", "elf32-bigarm", "elf32-littlearm")
OUTPUT_ARCH(arm)
ENTRY(main)

MEMORY
{
    ram (rwx) : ORIGIN = 0x20000000, LENGTH = 64k
}

SECTIONS
{
    .text :
    {
        . = ALIGN(4);
        *(.text*)
        . = ALIGN(4);
        _etext = .;
    } >ram
}
end = .;
```

# What does an executable image file contain?

- **.text segment**
  - Executable code
  - Initial reset vector
- **.data segment (.rodata in ELF)**
  - Static (initialized) variables
- **.bss segment**
  - Static (uninitialized) variables
  - Zero-filled by CRT or OS
  - From: Block Started by Symbol
- Does not contain heap or stack
- For details, see:  
`/usr/include/linux/elf.h`



# How can the contents of an executable file be read?

- Exactly the same way as an object file!
- Recall the useful options
  - -a show all information
  - -s shows symbols
  - -S shows section headers

# What does an executable file contain?

- Use readelf's -S option
- Note that the .out has fewer sections than the .o file
  - Why?

```
$ readelf -S example1.out
There are 6 section headers, starting at offset 0x8068:
```

## Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	00000000	000000	000000	00		0	0	0
[ 1]	.text	PROGBITS	00000000	008000	000014	00	AX	0	0	4
[ 2]	.ARM.attributes	ARM_ATTRIBUTES	00000000	008014	000021	00		0	0	1
[ 3]	.shstrtab	STRTAB	00000000	008035	000031	00		0	0	1
[ 4]	.symtab	SYMTAB	00000000	008158	000130	10		5	9	4
[ 5]	.strtab	STRTAB	00000000	008288	000085	00		0	0	1

## Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings)  
I (info), L (link order), G (group), x (unknown)  
O (extra OS processing required) o (OS specific), p (processor specific)

# What are the contents of an executable's .text segment?

- Use readelf's -x option to hex dump a section
- 1<sup>st</sup> column shows memory address
- 2<sup>nd</sup> through 5<sup>th</sup> columns show data
- The initial **SP** and **PC** values are visible
- The **executable opcodes** are also visible

```
$ readelf -x .text example1.out
```

Hex dump of section '.text':

0x00000000	00080020	09000000	0a200021	09180138	... .....	.!....8
0x00000010	fcd1fee7				....	

# What are the raw contents of an executable file?

- Use hexdump
- ELF's magic number is visible
- The initial SP, PC, executable opcodes are visible

```
$ hexdump example1.out
```

```
00000000 457f 464c 0101 0001 0000 0000 0000 0000  
00000010 0002 0028 0001 0000 0000 0000 0034 0000  
00000020 8068 0000 0000 0500 0034 0020 0001 0028  
00000030 0006 0003 0001 0000 8000 0000 0000 0000  
00000040 0000 0000 0014 0000 0014 0000 0005 0000  
00000050 8000 0000 0000 0000 0000 0000 0000 0000  
00000060 0000 0000 0000 0000 0000 0000 0000 0000  
*  
0008000 0800 2000 0009 0000 200a 2100 1809 3801  
0008010 d1fc e7fe 2041 0000 6100 6165 6962 0100  
0008020 0016 0000 4305 524f 4554 2d58 334d 0600  
0008030 070a 094d 0002 732e 6d79 6174 0062 732e  
0008040 7274 6174 0062 732e 7368 7274 6174 0062  
0008050 742e 7865 0074 412e 4d52 612e 7474 6972  
0008060 7562 6574 0073 0000 0000 0000 0000 0000  
0008070 0000 0000 0000 0000 0000 0000 0000 0000  
*  
0008090 001b 0000 0001 0000 0006 0000 0000 0000
```

# What purpose does an executable file serve?

- Serves as a convenient container for sections/segments
- Keeps segments segregate by type and access rights
- Serves as a program “image” for operating systems
- Allows the loader to place segments into main memory
- Can integrate symbol table and debugging information

How useful is an executable image  
for most embedded systems & tools?

# What does a binary program image contain?

- Basically, a binary copy of program's .text section
- Try ‘hexdump -C example.bin’
- Want to change the program?
  - Try ‘hexedit example.bin’
  - You can change the program (e.g. opcodes, static data, etc.)
- The initial SP, PC, executable opcodes are visible

```
$ hexdump -C example.bin
00000000  00 08 00 20 09 00 00 00  0a 20 00 21 09 18 01 38  |....|
00000010  fc d1 fe e7                      |....|
00000014
```

```
$ hexedit example.bin
00000000  00 08 00 20 09 00 00 00  0A 20 00 21 09 18 01 38  ....
00000010  FC D1 FE E7                      ....
```

# What are other, more usable formats?

- .o, .out, and .bin are all binary formats
- Many embedded tools don't use binary formats
- Two common ASCII formats
  - Intel hex (ihex)
  - Motorola S-records (srec)
- The initial **SP**, **PC**, **executable opcodes** are visible

```
$ arm-none-eabi-objcopy -O ihex example1.out "example1.hex "
```

```
$ cat example1.hex
```

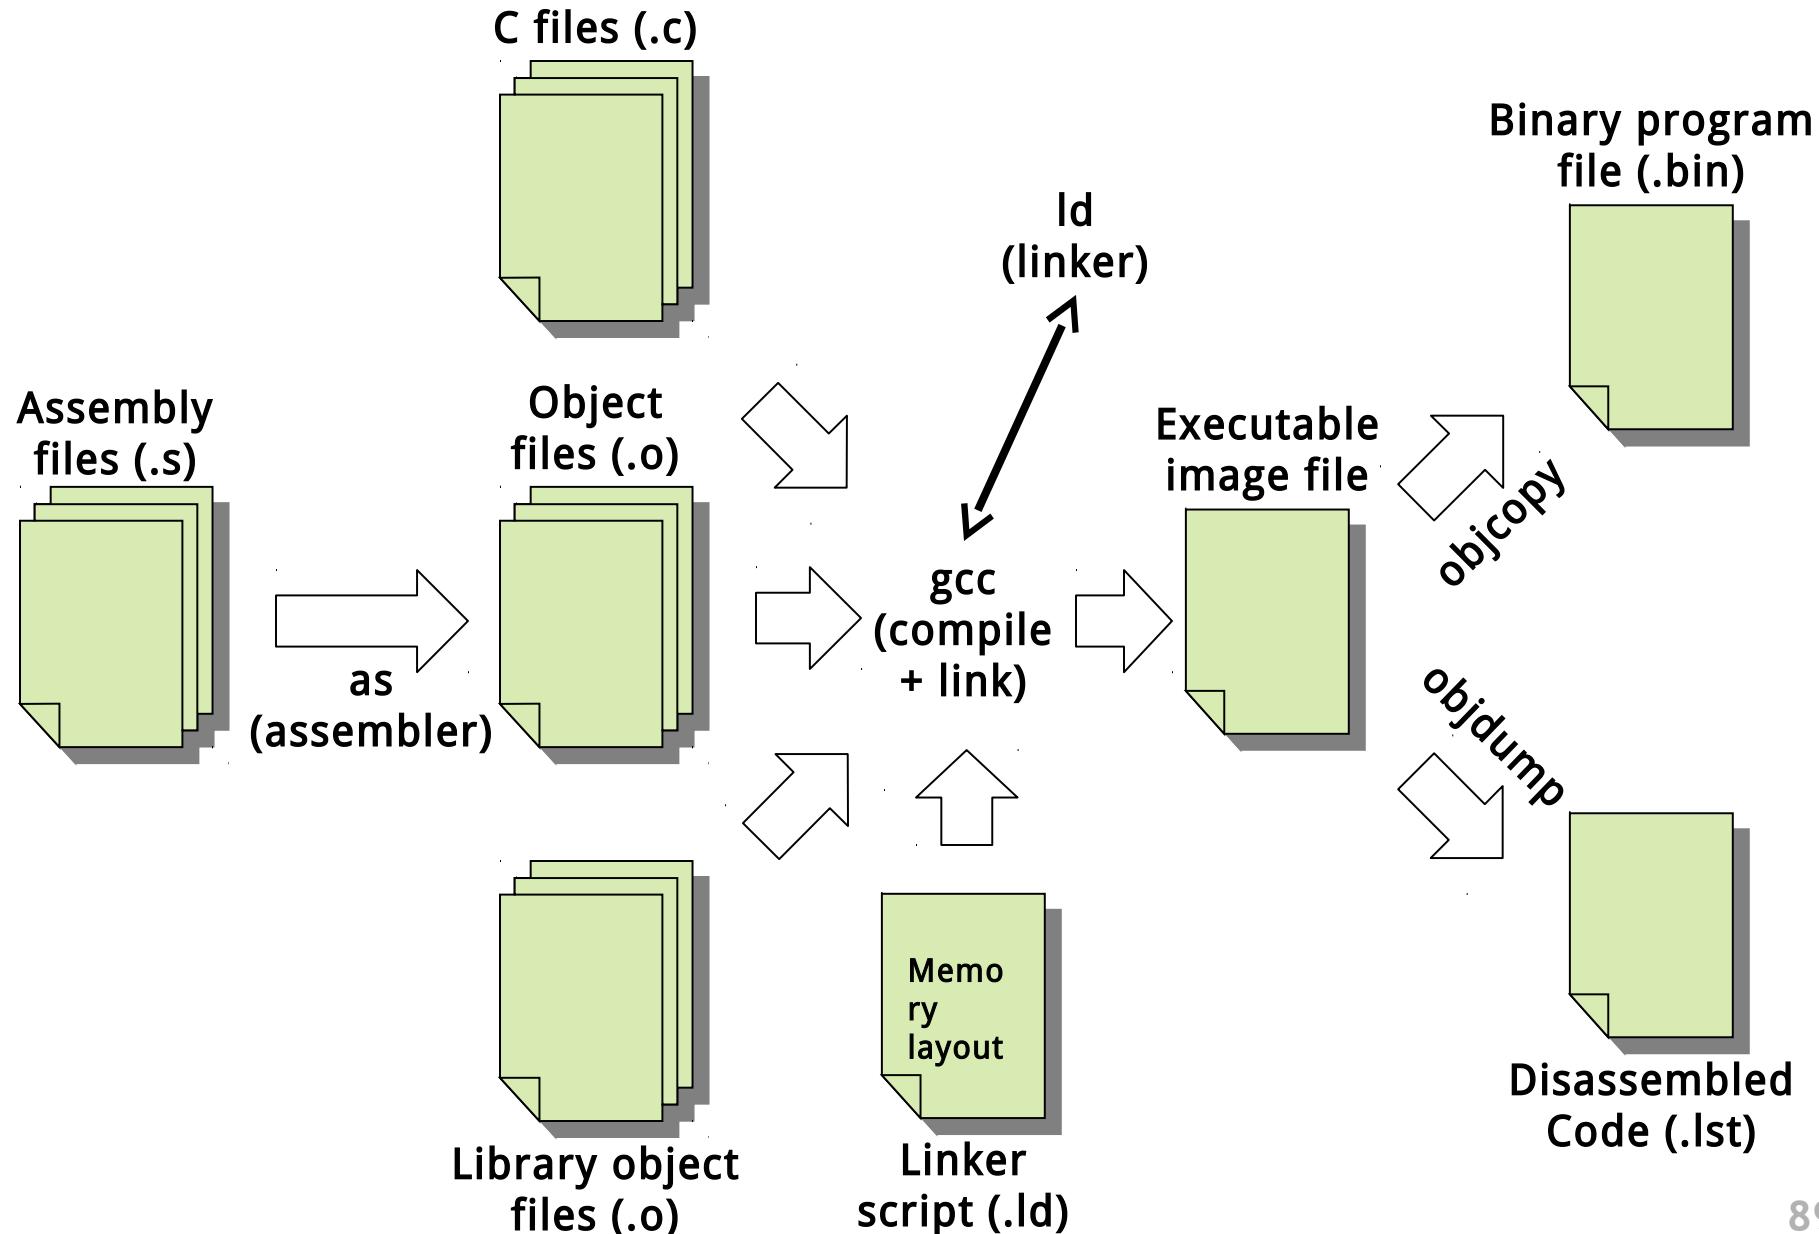
```
:100000000080020090000000A200021091801381A  
:04001000FCD1FEE73A  
:00000001FF
```

```
$ arm-none-eabi-objcopy -O srec example1.out "example1.srec "
```

```
$ cat example1.srec
```

```
S01000006578616D706C65312E73726563F7  
S113000000080020090000000A2000210918013816  
S1070010FCD1FEE736  
S9030000FC
```

# How does a mixed C/Assembly program get turned into a executable program image?



# generic hosted build

```
int factorial(int n) {
    int c;
    int result = 1;
    for (c = 1; c <= n; c++)
        result *= c;
    return result;
}
int main() {
    int i;
    int n;
    for (i = 0; i < 10; ++i) {
        n = factorial(i);
        printf("factorial(%d) = %d\n", i, n);
    }
}
```

```
$ arm-none-eabi-gcc \
-mcpu=cortex-m3 \
-mthumb main.c \
-T generic-hosted.ld \
-o factorial
$ qemu-arm -cpu cortex-m3 \
./factorial
factorial(0) = 1
factorial(1) = 1
factorial(2) = 2
factorial(3) = 6
factorial(4) = 24
factorial(5) = 120
factorial(6) = 720
factorial(7) = 5040
factorial(8) = 40320
factorial(9) = 362880
```

sudo apt-get install qemu-system qemu-user qemu-utils

# ARM Cortex-M3 Features

- Thumb-2 Instruction Set
- Bit Banding
- Integrated Peripherals
  - NVIC
  - Memory Protection Unit (MPU)
  - Debug Peripherals

# System Memory Map

Memory Map of Cortex-M3		Memory Map of FPGA Fabric Master, Ethernet MAC, Peripheral DMA
System Registers		0xE0043000 – 0xFFFF2FFF
External Memory Type 1	External Memory Type 1	0xE0042000 – 0xE0042FFF
External Memory Type 0	External Memory Type 0	0x78000000 – 0xE0041FFF
		0x74000000 – 0x77FFFFFF
		0x70000000 – 0x73FFFFFF
eNVM Controller	eNVM Controller	0x601D0000 – 0x6FFFFFFF
eNVM Aux Block (spare pages)	eNVM Aux Block (spare pages)	0x60180000 – 0x601CFFFF
eNVM Aux Block (array)	eNVM Aux Block (array)	0x60100100 – 0x6017FFFF
eNVM Spare Pages	eNVM Spare Pages	0x60100000 – 0x601000FF
eNVM Array	eNVM Array	0x60088200 – 0x600FFFFF
Peripherals (BB view)		0x60088000 – 0x600881FF
FPGA Fabric	FPGA Fabric	0x60084000 – 0x60087FFF
FPGA Fabric eSRAM Backdoor	FPGA Fabric eSRAM Backdoor	0x60080000 – 0x60083FFF
		0x60000000 – 0x6007FFFF
		0x44000000 – 0x5FFFFFFF
APB Extension Register		0x42000000 – 0x43FFFFFF
Analog Compute Engine	Analog Compute Engine	0x40100000 – 0x41FFFFFF
IAP Controller	IAP Controller	0x40050000 – 0x400FFFFF
eFROM	eFROM	0x40040000 – 0x4004FFFF
RTC	RTC	0x40030004 – 0x4003FFFF
MSS GPIO	MSS GPIO	0x40030000 – 0x40030003 → Visible only to FPGA Fabric Master
I2C_1	I2C_1	0x40017000 – 0x4001FFFFFF
SPI_1	SPI_1	0x40016000 – 0x40016FFF
UART_1	UART_1	0x40015000 – 0x40015FFF
		0x40014000 – 0x40014FFF
		0x40013000 – 0x40013FFF
		0x40012000 – 0x40012FFF
		0x40011000 – 0x40011FFF
		0x40010000 – 0x40010FFF
		0x40008000 – 0x4000FFFF
Fabric Interface Interrupt Controller	Fabric Interface Interrupt Controller	0x40007000 – 0x40007FFF
Watchdog	Watchdog	0x40006000 – 0x40006FFF
Timer	Timer	0x40005000 – 0x40005FFF
Peripheral DMA	Peripheral DMA	0x40004000 – 0x40004FFF
Ethernet MAC	Ethernet MAC	0x40003000 – 0x40003FFF
I2C_0	I2C_0	0x40002000 – 0x40002FFF
SPI_0	SPI_0	0x40001000 – 0x40001FFF
UART_0	UART_0	0x40000000 – 0x40000FFF
eSRAM_0 / eSRAM_1 (BB view)		0x24000000 – 0x3FFFFFFF
eSRAM_1	eSRAM_1	0x22000000 – 0x23FFFFFF
eSRAM_0	eSRAM_0	0x20010000 – 0x21FFFFFF
		0x20008000 – 0x2000FFFF
		0x20000000 – 0x20007FFF
		0x00088200 – 0x1FFFFFFF
eNVM (Cortex-M3) Virtual View	eNVM (fabric) Virtual View	0x000881FF } Visible only to FPGA Fabric Master
		0x00000000

Figure 2-4 • System Memory Map with 64 Kbytes of SRAM

# Memory-mapped I/O

- The idea is really simple
  - Instead of real memory at a given memory address, have an I/O device respond.
- Example:
  - Let's say we want to have an LED turn on if we write a "1" to memory location 5.
  - Further, let's have a button we can read (pushed or unpushed) by reading address 4.
    - If pushed, it returns a 1.
    - If not pushed, it returns a 0.

# Accessing memory locations from C

- Memory has an address and value
- Can equate a pointer to desired address
- Can set/get de-referenced value to change memory

```
#define SYSREG_SOFT_RST_CR 0xE0042030

uint32_t *reg = (uint32_t *)(SYSREG_SOFT_RST_CR);

main () {
    *reg |= 0x00004000; // Reset GPIO hardware
    *reg &= ~(0x00004000);
}
```

# Some useful C keywords

- **const**
  - Makes variable value or pointer parameter unmodifiable
  - `const foo = 32;`
- **register**
  - Tells compiler to locate variables in a CPU register if possible
  - Useless in C99
  - `register int x;`
- **static**
  - Preserve variable value after its scope ends
  - Does not go on the stack
  - `static int x;`
- **volatile**
  - Opposite of const
  - Can be changed in the background
  - `volatile int l;`

# What happens when **this** “instruction” executes?

```
#include <stdio.h>
#include <inttypes.h>

#define REG_FOO 0x40000140

main () {
    uint32_t *reg = (uint32_t *)(REG_FOO);
    *reg += 3;

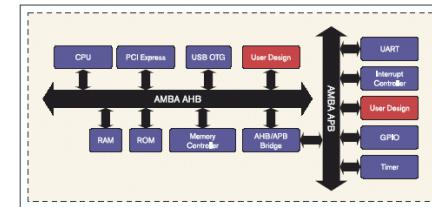
    printf( “0x%x\n” , *reg);
}
```

# **“\*reg += 3” is turned into a ld, add, str sequence**

- Load instruction
  - A bus read operation commences
  - The CPU drives the address “reg” onto the address bus
  - The CPU indicated a read operation is in process (e.g. R/W#)
  - Some “handshaking” occurs
  - The target drives the contents of “reg” onto the data lines
  - The contents of “reg” is loaded into a CPU register (e.g. r0)
- Add instruction
  - An immediate add (e.g. add r0, #3) adds three to this value
- Store instruction
  - A bus write operation commences
  - The CPU drives the address “reg” onto the address bus
  - The CPU indicated a write operation is in process (e.g. R/W#)
  - Some “handshaking” occurs
  - The CPU drives the contents of “r0” onto the data lines
  - The target stores the data value into address “reg”

# Details of the bus “handshaking” depend on the particular memory/peripherals involved

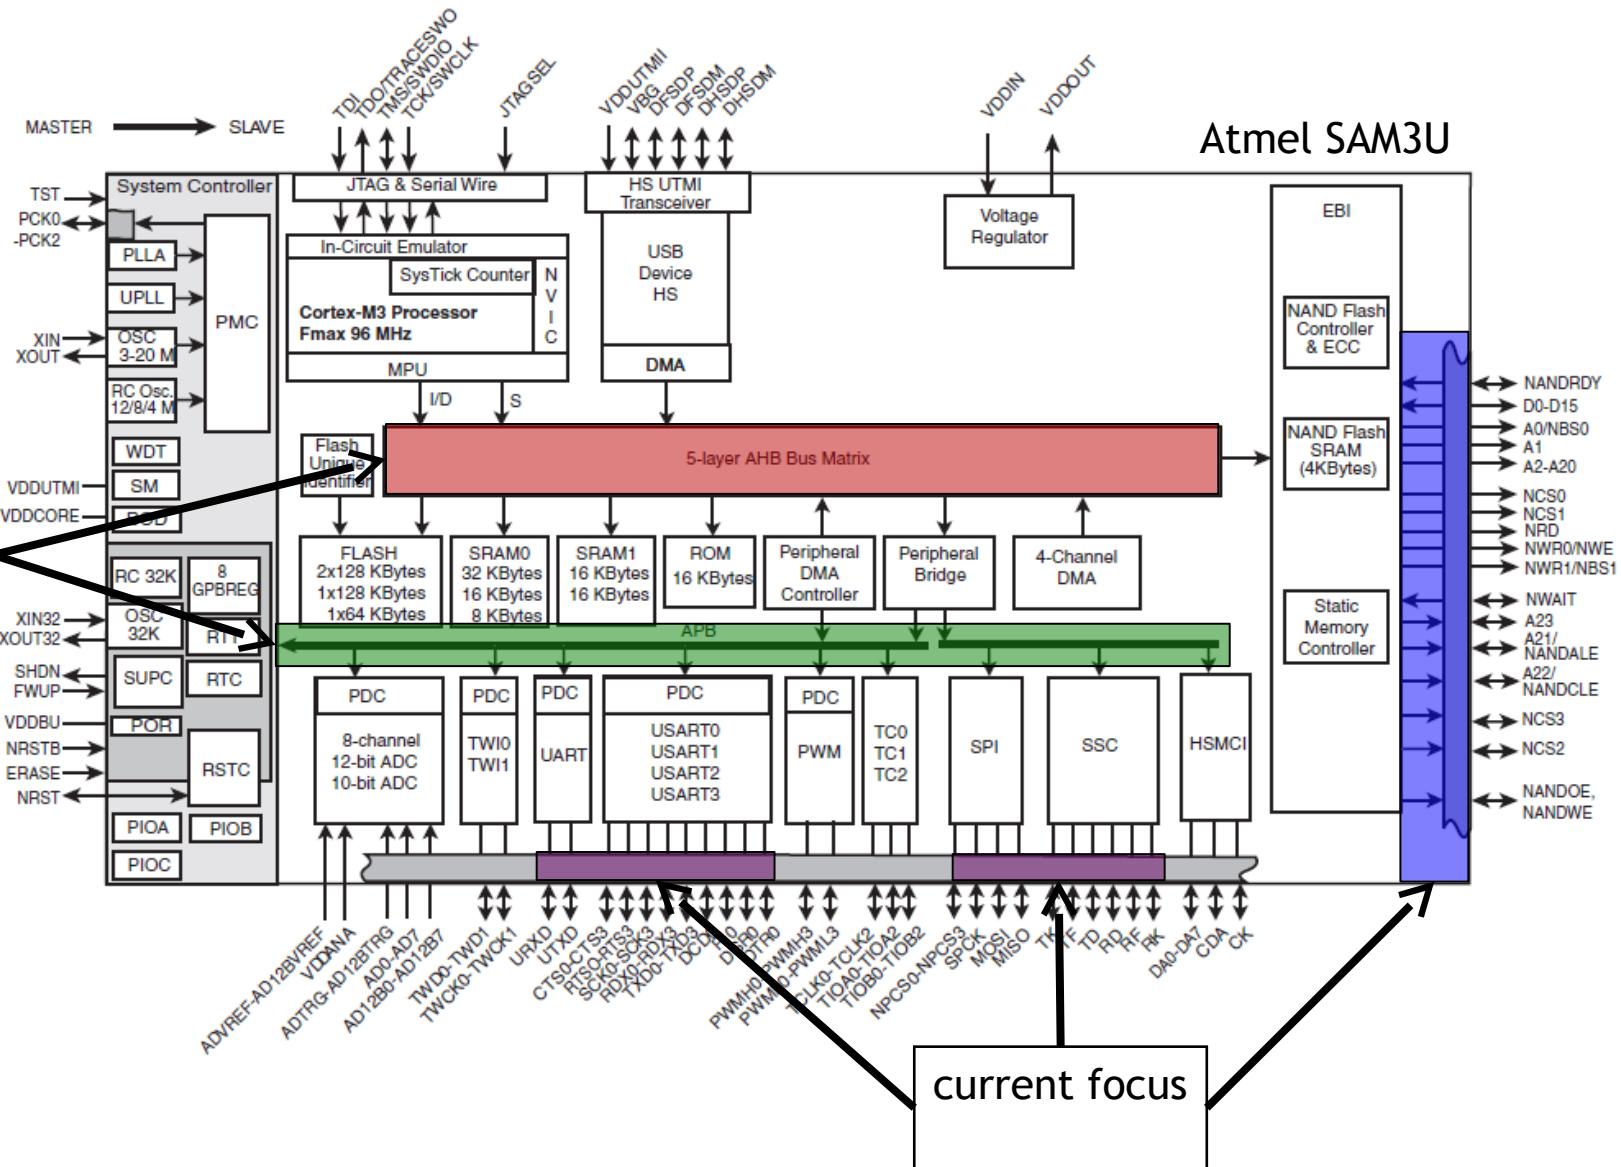
- SoC memory/peripherals
  - AMBA AHB/APB



- NAND Flash
  - Open NAND Flash Interface (ONFI)
- DDR SDRAM
  - JEDEC JESD79, JESD79-2F, etc.



# Modern embedded systems have multiple busses

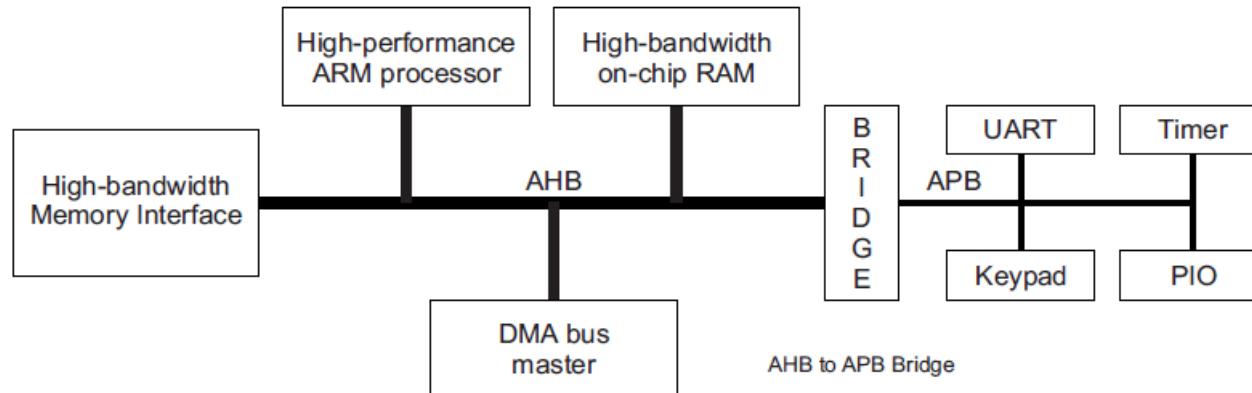


# Why have so many busses?

- Many designs considerations
  - Master vs Slave
  - Internal vs External
  - Bridged vs Flat
  - Memory vs Peripheral
  - Synchronous vs Asynchronous
  - High-speed vs low-speed
  - Serial vs Parallel
  - Single master vs multi master
  - Single layer vs multi layer
  - Multiplexed A/D vs demultiplexed A/D
- Discussion: what are some of the tradeoffs?

# Advanced Microcontroller Bus Architecture (AMBA)

- Advanced High-performance Bus (AHB)
- Advanced Peripheral Bus (APB)



## AHB

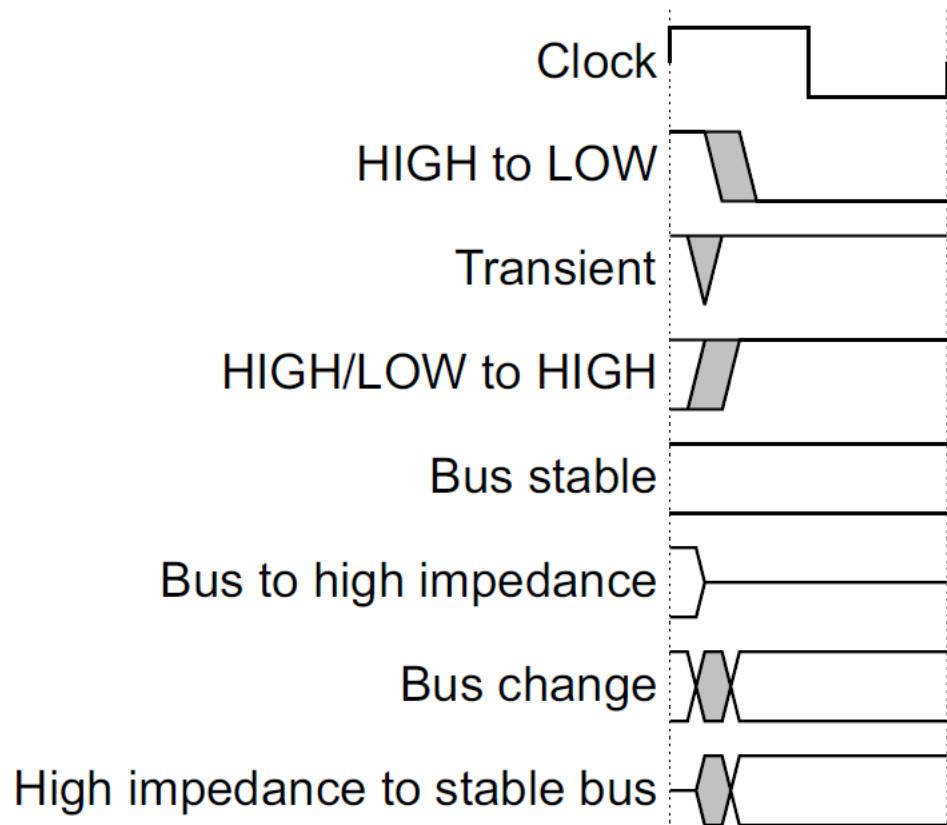
- High performance
- Pipelined operation
- Burst transfers
- Multiple bus masters
- Split transactions

## APB

- Low power
- Latched address/control
- Simple interface
- Suitable of many peripherals

# Key to timing diagram conventions

- Timing diagrams
  - Clock
  - Stable values
  - Transitions
  - High-impedance



- Signal conventions

- Lower case 'n' denote active low (e.g. RESETn)
- Prefix 'H' denotes AHB
- Prefix 'P' denotes APB

# Basic read and write transfers with no wait states

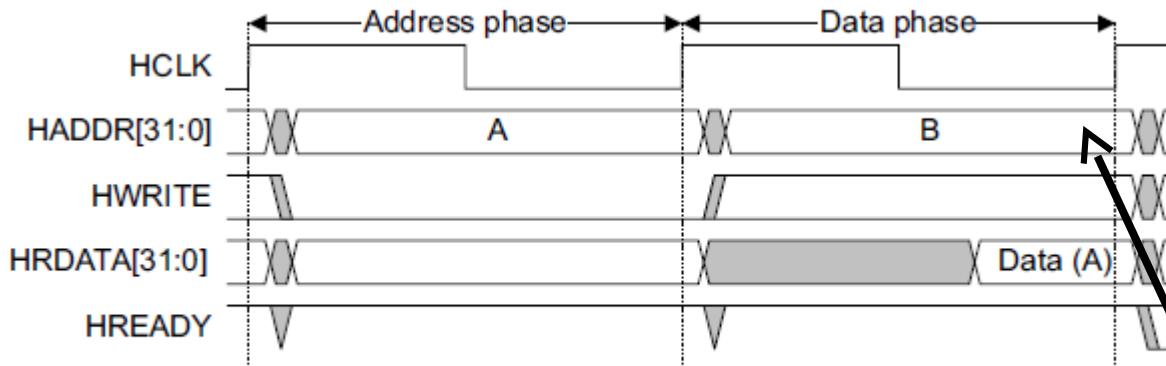


Figure 3-1 Read transfer

Pipelined  
Address  
& Data  
Transfer

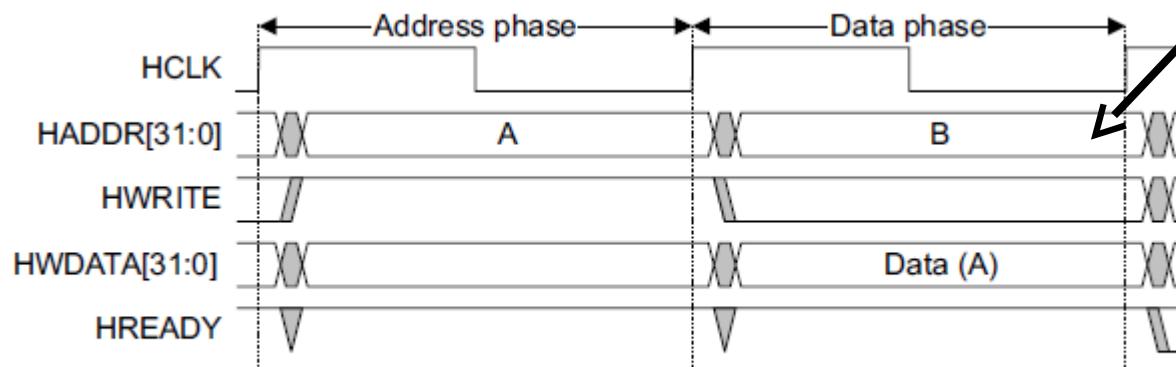
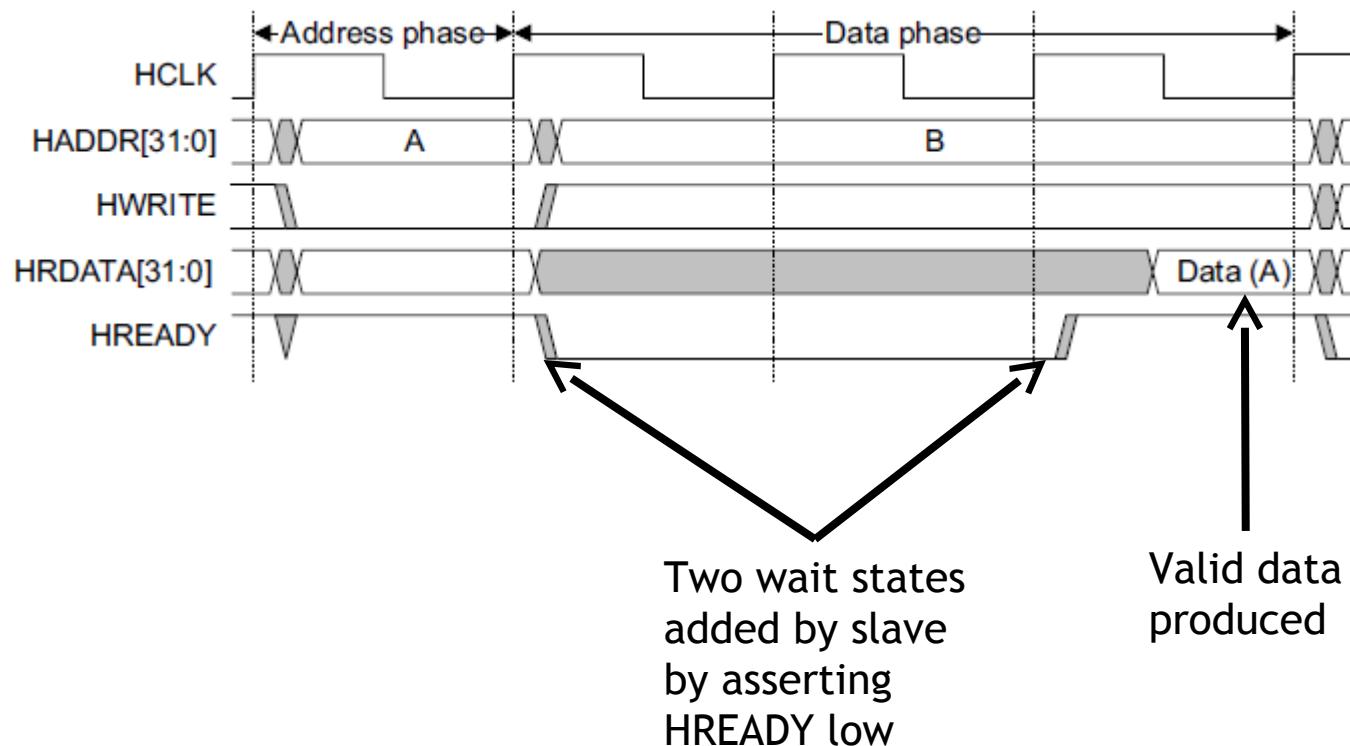
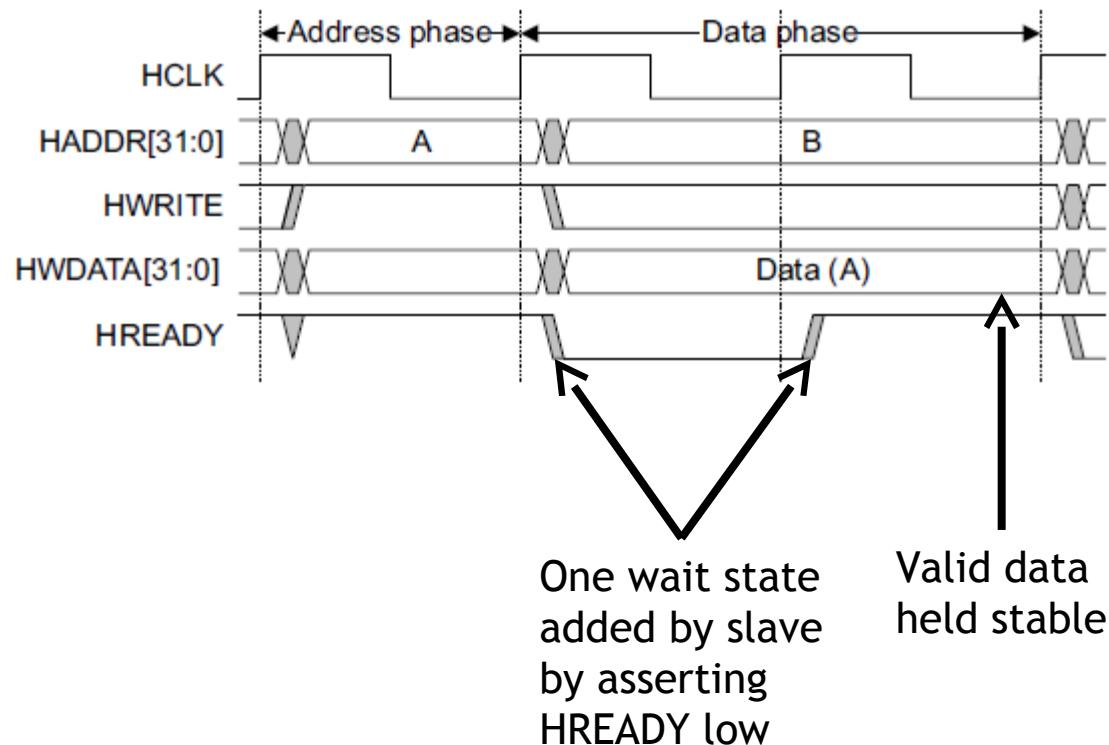


Figure 3-2 Write transfer

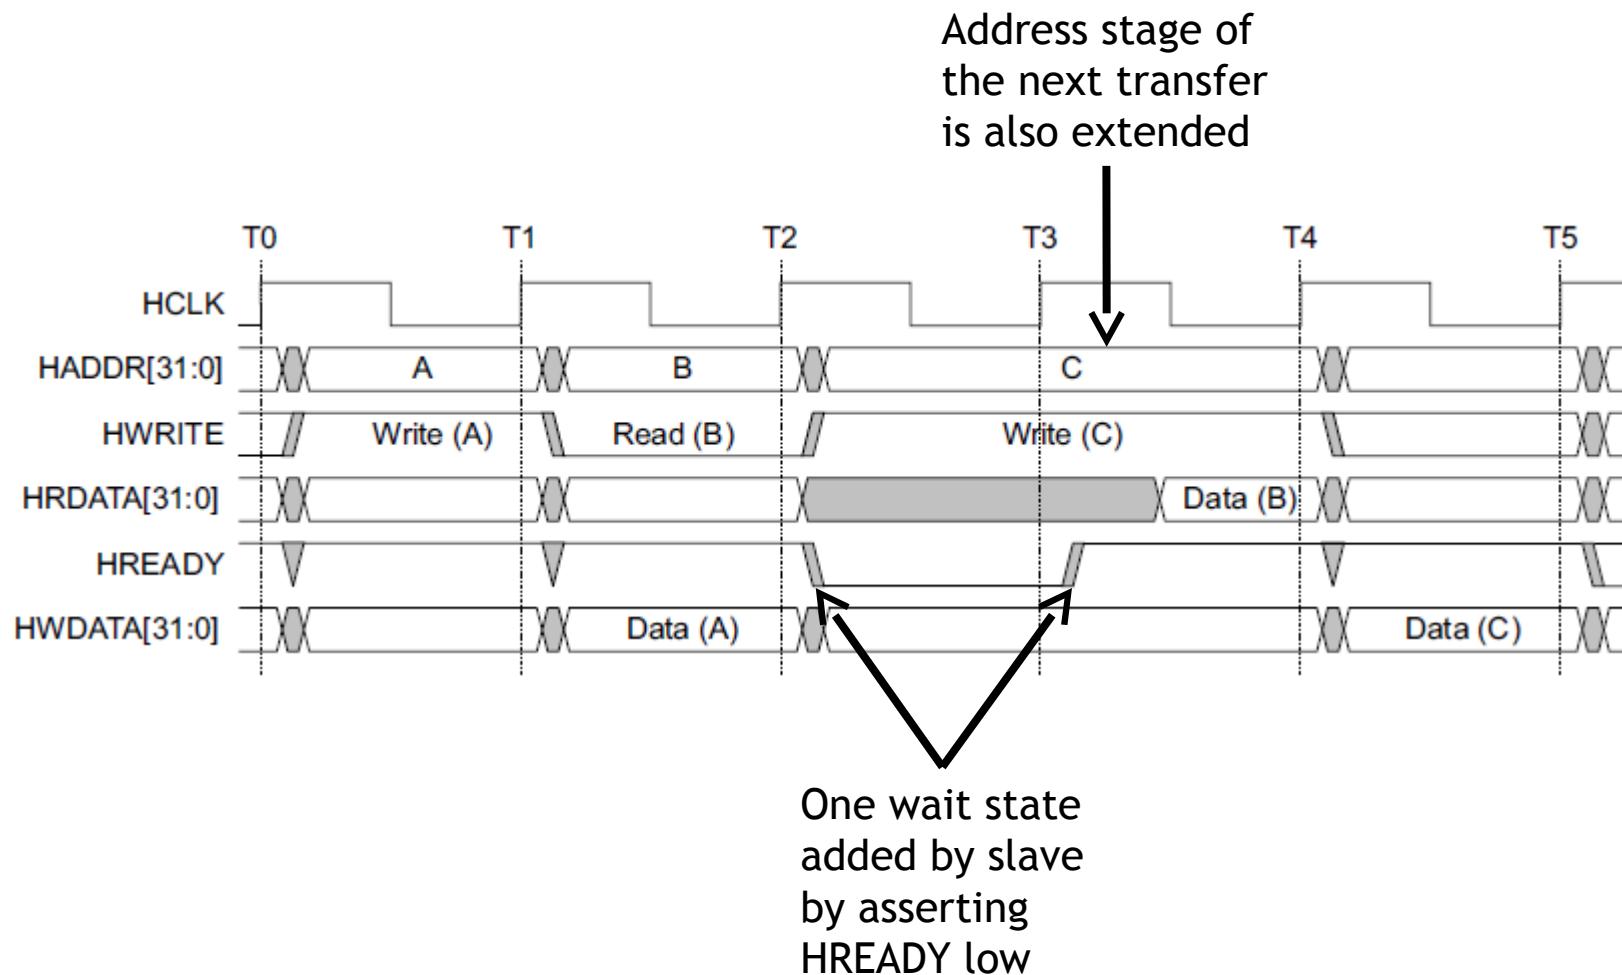
# Read transfer with two wait states



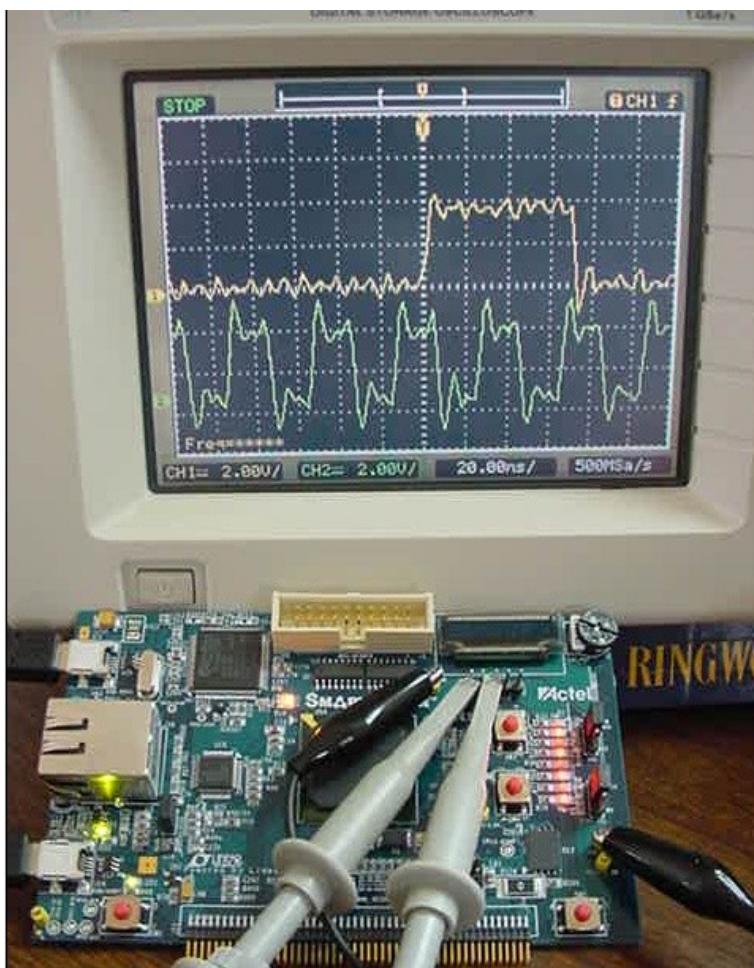
# Write transfer with one wait state



# Wait states extend the address phase of next transfer



# memory-mapped peripheral



# I/O Data Transfer

Two key questions to determine how data is transferred to/from a non-trivial I/O device:

1. How does the CPU know when data is available?
  - a. Polling
  - b. Interrupts
2. How is data transferred into and out of the device?
  - a. Programmed I/O
  - b. Direct Memory Access (DMA)

## Two basic types of interrupts

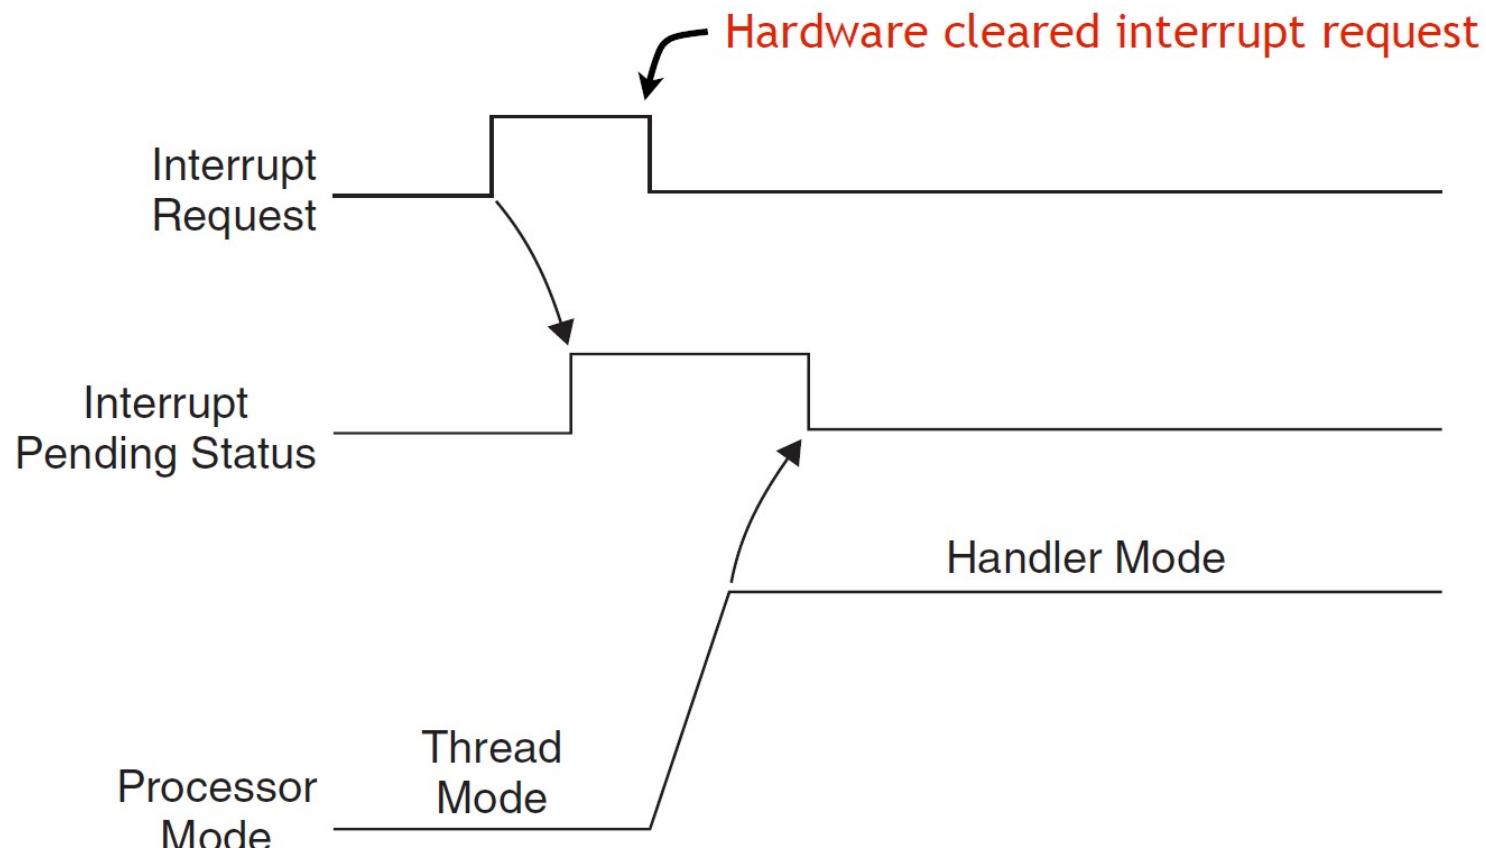
(1/2)

- Those caused by an instruction
  - Examples:
    - TLB miss
    - Illegal/unimplemented instruction
    - div by 0
  - Names:
    - Trap, exception

## Two basic types of interrupts (2/2)

- Those caused by the external world
  - External device
  - Reset button
  - Timer expires
  - Power failure
  - System error
- Names:
  - interrupt, external interrupt

# Pending interrupts



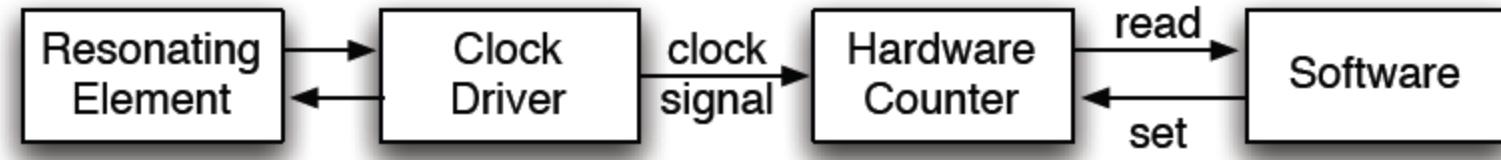
The normal case. Once Interrupt request is seen, processor puts it in “pending” state even if hardware drops the request.  
IPS is cleared by the hardware once we jump to the ISR.

## Fine grain motion control

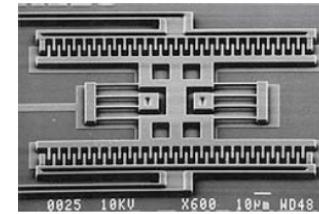
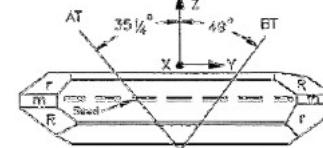


<http://www.youtube.com/watch?v=SOESSCXGhFo>

# Clock generation and use



- Resonating element/Driver:
  - Quartz crystal can be made to resonate due to Piezoelectric effect.
    - Resonate frequency depends on length, thickness, and angle of cut.
    - Issues: Very stable (<100ppm) but not all frequencies possible.
  - MEMS Resonator
    - Arbitrary frequency, potentially cheap, susceptible to temperature variations.
  - Others:
    - Inverter Ring, LC/RC circuits, Atomic clock, and many more.



# Appendix

# Branch

Table A4-1 Branch instructions

Instruction	Usage	Range
<a href="#">B</a> on page A6-40	Branch to target address	+/-1 MB
<a href="#">CBNZ</a> , <a href="#">CBZ</a> on page A6-52	Compare and Branch on Nonzero, Compare and Branch on Zero	0-126 B
<a href="#">BL</a> on page A6-49	Call a subroutine	+/-16 MB
<a href="#">BLX (register)</a> on page A6-50	Call a subroutine, optionally change instruction set	Any
<a href="#">BX</a> on page A6-51	Branch to target address, change instruction set	Any
<a href="#">TBB</a> , <a href="#">TBH</a> on page A6-258	Table Branch (byte offsets)	0-510 B
	Table Branch (halfword offsets)	0-131070 B

# Data processing instructions

Table A4-2 Standard data-processing instructions

Mnemonic	Instruction	Notes
ADC	Add with Carry	-
ADD	Add	Thumb permits use of a modified immediate constant or a zero-extended 12-bit immediate constant.
ADR	Form PC-relative Address	First operand is the PC. Second operand is an immediate constant. Thumb supports a zero-extended 12-bit immediate constant. Operation is an addition or a subtraction.
AND	Bitwise AND	-
BIC	Bitwise Bit Clear	-
CBN	Compare Negative	Sets flags. Like ADD but with no destination register.
CMP	Compare	Sets flags. Like SUB but with no destination register.
EOR	Bitwise Exclusive OR	-
MOV	Copies operand to destination	Has only one operand, with the same options as the second operand in most of these instructions. If the operand is a shifted register, the instruction is an LSL, LSR, ASR, or ROR instruction instead. See <i>Shift instructions</i> on page A4-10 for details. Thumb permits use of a modified immediate constant or a zero-extended 16-bit immediate constant.

Many, Many More!

# Load/Store instructions

Table A4-10 Load and store instructions

Data type	Load	Store	Load unprivileged	Store unprivileged	Load exclusive	Store exclusive
32-bit word	LDR	STR	LDRT	STRT	LDREX	STREX
16-bit halfword	-	STRH	-	STRHT	-	STREXH
16-bit unsigned halfword	LDRH	-	LDRHT	-	LDREXH	-
16-bit signed halfword	LDRSH	-	LDRSHT	-	-	-
8-bit byte	-	STRB	-	STRBT	-	STREXB
8-bit unsigned byte	LDRB	-	LDRBT	-	LDREXB	-
8-bit signed byte	LDRSB	-	LDRSBT	-	-	-
two 32-bit words	LDRD	STRD	-	-	-	-

# Miscellaneous instructions

Table A4-12 Miscellaneous instructions

Instruction	See
Clear Exclusive	<i>CLREX</i> on page A6-56
Debug hint	<i>DBG</i> on page A6-67
Data Memory Barrier	<i>DMB</i> on page A6-68
Data Synchronization Barrier	<i>DSB</i> on page A6-70
Instruction Synchronization Barrier	<i>ISB</i> on page A6-76
If Then (makes following instructions conditional)	<i>IT</i> on page A6-78
No Operation	<i>NOP</i> on page A6-167
Preload Data	<i>PLD</i> , <i>PLDW</i> ( <i>immediate</i> ) on page A6-176 <i>PLD</i> ( <i>register</i> ) on page A6-180
Preload Instruction	<i>PLI</i> ( <i>immediate, literal</i> ) on page A6-182 <i>PLI</i> ( <i>register</i> ) on page A6-184
Send Event	<i>SEV</i> on page A6-212
Supervisor Call	<i>SVC</i> ( <i>formerly SWT</i> ) on page A6-252
Wait for Event	<i>WFE</i> on page A6-276
Wait for Interrupt	<i>WFI</i> on page A6-277
Yield	<i>YIELD</i> on page A6-278

# Conditional execution: Append to many instructions for conditional execution

Table A6-1 Condition codes

cond	Mnemonic extension	Meaning (integer)	Meaning (floating-point) <sup>ab</sup>	Condition flags
0000	EQ	Equal	Equal	Z = 1
0001	NE	Not equal	Not equal, or unordered	Z = 0
0010	CS <sup>c</sup>	Carry set	Greater than, equal, or unordered	C = 1
0011	CC <sup>d</sup>	Carry clear	Less than	C = 0
0100	MI	Minus, negative	Less than	N = 1
0101	PL	Plus, positive or zero	Greater than, equal, or unordered	N = 0
0110	VS	Overflow	Unordered	V = 1
0111	VC	No overflow	Not unordered	V = 0
1000	HI	Unsigned higher	Greater than, or unordered	C = 1 and Z = 0
1001	LS	Unsigned lower or same	Less than or equal	C = 0 or Z = 1
1010	GE	Signed greater than or equal	Greater than or equal	N = V
1011	LT	Signed less than	Less than, or unordered	N ≠ V
1100	GT	Signed greater than	Greater than	Z = 0 and N = V
1101	LE	Signed less than or equal	Less than, equal, or unordered	Z = 1 or N ≠ V
1110	None (AL) <sup>e</sup>	Always (unconditional)	Always (unconditional)	Any